

# Lecture 8: Declarative Goals and Planning in BDI Agent-Oriented Programming Languages

Autonomous Agents and Multiagent Systems  
DIS, La Sapienza - PhD Course

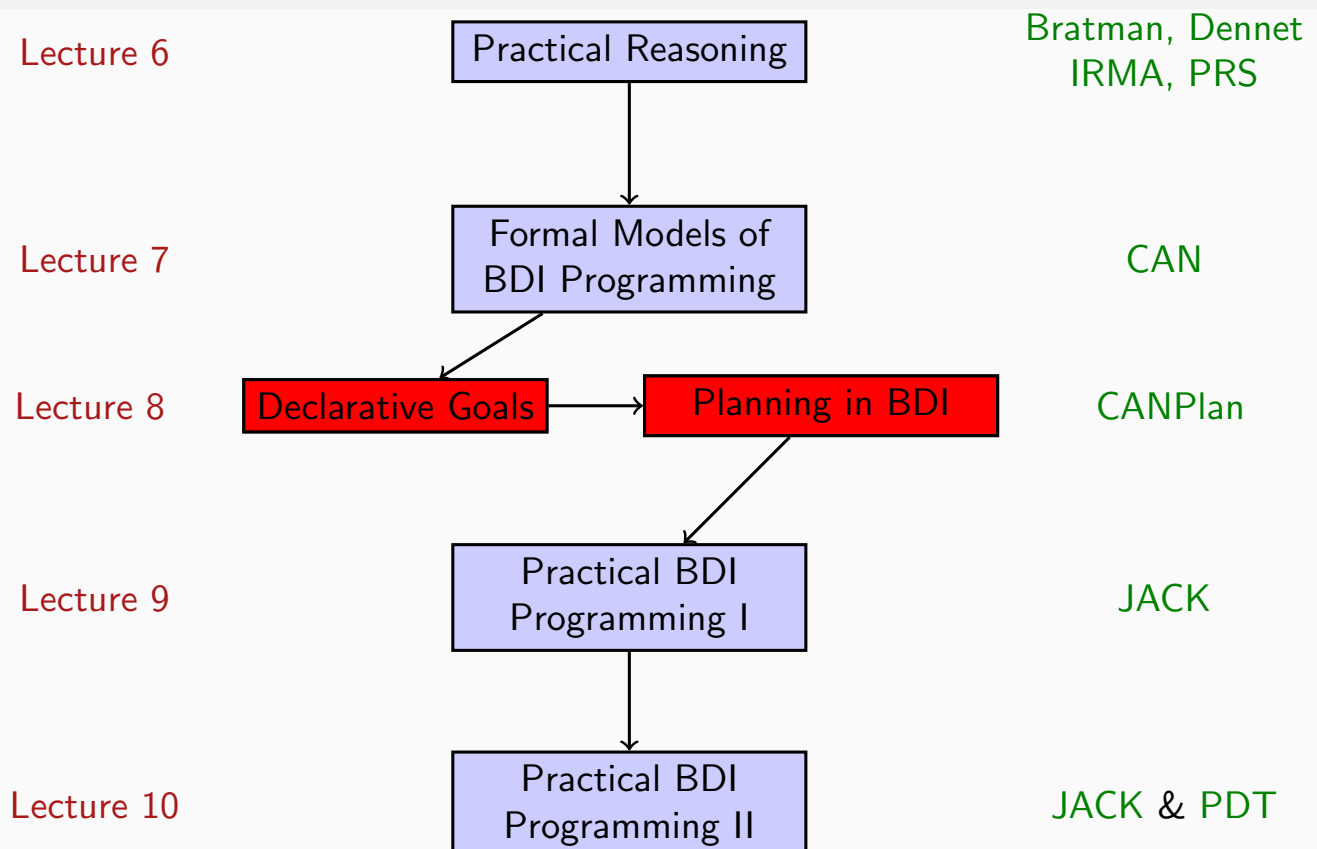
Sebastian Sardina<sup>1</sup>

<sup>1</sup>Department of Computer Science and Information Technology  
RMIT University  
Melbourne, AUSTRALIA



November 29, 2007

## Roadmap for Next Lectures



In the previous lecture “*Formal Models of BDI Programming*” we have seen:

**1 Basic concepts of BDI programming:**

- ▶ programming using mentalistic concepts such as beliefs, desires, capabilities, etc.
- ▶ goal-oriented programming – via **events**;
- ▶ implicit programming – via **plan library & context conditions**;
- ▶ rational execution cycle: on-the-fly recombination of plans.

**2 CAN formal BDI programming language:**

- ▶ captures the basic notions of BDI programming: rational executor;
- ▶ formal operational semantics;
- ▶ includes built-in failure handling.

## This Lecture: Declarative Goals & Planning

In the next lecture we show how to:

**1 Accommodate declarative goals into CAN.**

- ▶ to provide an hybrid account of goals with both declarative & procedural aspects.

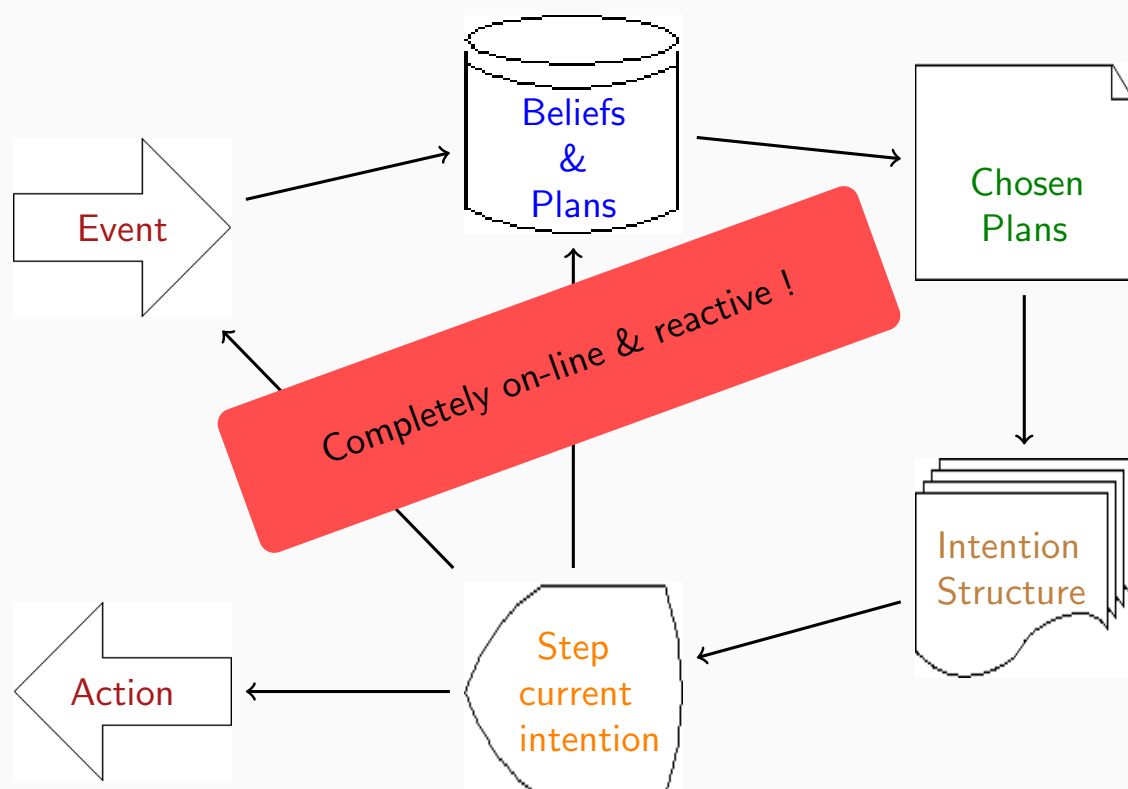
**2 Accommodate hierarchical HTN-style planning into CAN.**

- ▶ to perform some “offline” lookahead reasoning within the whole online “reactive” execution scheme.

# Outline

- 1 Review of CAN
- 2 Declarative Goals in CAN
- 3 CANPlan: CAN + HTN Planning
  - Motivation
  - HTN Planning
  - Plan Construct
- 4 Conclusions

## The BDI Execution Cycle [Rao&Georgeff 92]



# The CAN Language [Winikioff et al. 2002]

## CAN: Conceptual Agent Notation

Can be seen as an extension of Rao's AgentSpeak.

A CAN agent is defined as  $Agt = \langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle$ , where:

- ▶  $\mathcal{N}$  is the **agent name**.
- ▶  $\mathcal{B}$  is the **belief base**: current agent's knowledge.
- ▶  $\mathcal{A}$  is the sequence of actions executed so far.
- ▶  $\Pi$  is a **plan library** containing plan rules  $e : \psi \leftarrow P$ :
  - ▶  $e$  is the **triggering event**
  - ▶  $\psi$  is the **context condition**
  - ▶  $P$  is the **plan-body**
- ▶  $\Gamma$  is the **intention base**: partially uninstantiated plan-bodies.

## The CAN Language: Plans & Intentions

- ▶  $\Pi$  is a **plan library** containing plan rules  $e : \psi \leftarrow P$ :

$act$	<i>primitive action</i>
$+b$	<i>belief addition</i>
$-b$	<i>belief deletion</i>
$? \phi$	<i>tests goal</i>
$!e$	<i>posting of achievement event goal</i>
$P_1; P_2$	<i>sequence</i>
$P_1    P_2$	<i>interleaved concurrency</i>

Plus, the following system-auxiliary constructs:

$nil(\theta)$	<i>empty program with bindings</i>
$P_1 \triangleright P_2$	<i>try <math>P_1</math>; else <math>P_2</math></i>
$(\psi_1 : P_1, \dots, \psi_n : P_n)$	<i>guarded plans</i>

- ▶  $\Gamma$  is the **intention base**: set of partially uninstantiated plan-bodies.
  - ▶ E.g.:  $(?phone(john, N); call(N); !talk) || !cook\_dinner$

# Semantics of CAN

The semantics of CAN is modularly defined in two levels:

- 1 Agent-level semantics:  $\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', \Gamma' \rangle$ .
  - State that agent configuration  $\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle$  **may legally evolve** to configuration  $\langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', \Gamma' \rangle$ .
- 2 Intention-level semantics:  $\langle \Pi, \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \langle \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$ .
  - State that intention configuration  $\langle \Pi, \mathcal{B}, \mathcal{A}, P \rangle$  **may legally evolve** to configuration  $\langle \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$ .

Legal transitions are characterized by a set of rules of the form:

$$\frac{\text{Set of conditions}}{C \longrightarrow C'} \text{ RuleName}$$

## Definition (BDI Agent Execution)

A BDI execution  $E$  of an agent  $C_0 = \langle \mathcal{N}, \Pi, \mathcal{B}_0, \mathcal{A}_0, \Gamma_0 \rangle$  is a, possibly infinite, sequence of agent configurations  $C_0 \cdot C_1 \cdot \dots$  such that  $C_i \Longrightarrow C_{i+1}$ , for every  $i \geq 0$ . A terminating execution is a finite execution  $C_0 \cdot \dots \cdot C_n$  with  $\Gamma_n = \{\}$ .

# Agent-Level Semantics

Assume  $\langle \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle$  is given.

$$\frac{P \in \Gamma \quad \langle \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', (\Gamma \setminus \{P\}) \cup \{P'\} \rangle} \text{Agt}_{\text{step}}$$

**Execute an active intention  $P$ .**

$$\frac{e \text{ is a new external event}}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \cup \{!e\} \rangle} \text{Agt}_{\text{event}}$$

**Assimilate an external event  $e$ .**

$$\frac{P \in \Gamma \quad \langle \mathcal{B}, \mathcal{A}, P \rangle \not\rightarrow}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \setminus \{P\} \rangle} \text{Agt}_{\text{clean}}$$

**Remove an active intention  $P$  that is blocked.**

# Intention-Level Semantics

$$\frac{\mathcal{B} \models \phi\theta}{\langle \mathcal{B}, \mathcal{A}, ?\phi \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, \text{nil}(\theta) \rangle} \text{?} \quad \frac{}{\langle \mathcal{B}, \mathcal{A}, \text{act} \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A} \cdot \text{act}, \text{nil}(\emptyset) \rangle} \text{do}$$

$$\frac{}{\langle \mathcal{B}, \mathcal{A}, +b \rangle \longrightarrow \langle \mathcal{B} \cup \{b\}, \mathcal{A}, \text{nil}(\emptyset) \rangle} +b \quad \frac{}{\langle \mathcal{B}, \mathcal{A}, -b \rangle \longrightarrow \langle \mathcal{B} \setminus \{b\}, \mathcal{A}, \text{nil}(\emptyset) \rangle} -b$$

$$\frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_1 \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1; P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_1; P_2 \rangle} \text{Seq} \quad \frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \parallel P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \parallel P_2 \rangle} \parallel_1$$

# Intention-Level Semantics: Selection & Failure

$$!e \longrightarrow \langle \psi_1 : P_1, \dots, \psi_n : P_n \rangle \longrightarrow$$

$$P_i \theta_i \triangleright \langle \psi_1 : P_1, \dots, \psi_i \wedge \vec{x} \neq \theta_i : P_i, \dots, \psi_n : P_n \rangle \xrightarrow{*}$$

$$\frac{\Delta = \{\psi_i \theta : P_i \theta \mid e' : \psi_i \leftarrow P_i \in \Pi \wedge \theta = \text{mgu}(e, e')\}}{\langle \mathcal{B}, \mathcal{A}, !e \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, \langle \Delta \rangle \rangle} \text{Event}$$

$$\frac{\psi_i(\vec{x}) : P_i \in \Delta \quad \mathcal{B} \models \psi_i(\vec{x})\theta}{\langle \mathcal{B}, \mathcal{A}, \langle \Delta \rangle \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, P_i \theta \triangleright \langle (\Delta \setminus \{\psi_i(\vec{x}) : P_i\}) \cup \{\psi_i(\vec{x}) \wedge \vec{x} \neq \theta : P_i\} \rangle \rangle} \text{Sel}$$

$$\frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \triangleright P_2 \rangle} \triangleright \quad \frac{}{\langle \mathcal{B}, \mathcal{A}, (\text{nil} \triangleright P_2) \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, \text{nil} \rangle} \triangleright_t$$

$$\frac{P_1 \neq \text{nil} \quad \langle \mathcal{B}, \mathcal{A}, P_1 \rangle \not\longrightarrow \langle \mathcal{B}, \mathcal{A}, P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle} \triangleright_f$$

## Goals-to-be vs Goals-to-do

Contrast the following two kind of goals:

**Goals-to-be** Bring about a state of affairs.

- ▶ Used in automated planning and agent theory.
- ▶ Have a **declarative** flavor.
- ▶ E.g., achieve *BeAtLocation(20, 12)*.

**Goals-to-do** Complete a task or procedure.

- ▶ Used in hierarchical planning and high-level programming (ConGolog).
- ▶ Have a **procedural** flavor.
- ▶ E.g., complete *goToLocation(20, 12)*.

Agent goals in BDI programming as modelled with **events**!

Agent goals in BDI programming as events are pure **goals-to-do**!

## Advantages of Declarative Goals

- 1 Help decouple plan execution and goal achievement.
  - ▶ Completing the plan **vs** achieving the goal.
- 2 Facilitate goal dynamics.
  - ▶ When is a (new) sub-goal required for another goal?
  - ▶ When should a goal be dropped? (e.g., it can never be obtained)
- 3 Facilitate plan failure handling
  - ▶ What can be done if a plan for a goal fails?
  - ▶ Is that the same as failing the goal itself?
- 4 Enable reasoning about goal and plan interaction
  - ▶ Does a plan *guarantee* a goal?
  - ▶ Can a plan *achieve* a goal?
  - ▶ Can a plan *preclude* a goal?
- 5 Enhance goal and plan communication
  - ▶ Communicate desires/requests **vs** communicate know-how.

## Towards a Hybrid Goals

We want to keep the procedural advantages of event-goals:

- ▶ Built-in know-how available from domain experts.
- ▶ Reduced search space to cope online with dynamic environments.
- ▶ BDI execution engine with failure handling.

However, we would like to bring in the following features of declarative goals:

**Declarative** To facilitate potential reasoning and communication.

**Persistent** A rational agent should not abandon a goal without good reasons.

**Unachieved** A rational agent should not be pursuing goals that are already true.

**Possible** A rational agent should only pursue goals that are eventually possible to achieve.

## Adding Declarative Goals to CAN

Normal event goals-to-do  $!e$  are extended with declarative information:

$\text{Goal}(\phi_s, !e, \phi_f)$

**Achieve** goal  $\phi_s$  by **using/posting** event  $e$ ; **failing** if  $\phi_f$  becomes true.

- 1  $\phi_s$  is the **success condition** – when the goal assumed achieved.
- 2  $e$  is the **know-how** of the goal – how the goal can be realized.
- 3  $\phi_f$  is the **failure condition** – when the goal is impossible or not needed.

### Example

$\text{Goal}(\neg \text{Hungry}, !\text{eatFood}, \neg \text{foodAvailable})$ .

OBS.: there may be many plans in the library for event *eatFood*.



# Goal-programs Properties

## Goal( $\phi_s, P, \phi_f$ ) Properties

- 1 If  $\phi_s$  becomes true, terminate successfully. [unachieved]
- 2 If  $\phi_f$  becomes true, terminate with failure. [possible]
- 3 If  $P$  ends successfully and  $\phi_s$  is still *not* true,  $P$  is re-tried. [persistent]
- 4 If  $P$  fails and  $\phi_s$  is still *not* true,  $P$  is re-tried. [persistent]

## Example (Goal( $\neg Hungry, !eatFood, \neg foodAvailable$ ))

- 1 If agent happens not to be hungry anymore (e.g., maybe feels sick), then the goal terminates successfully.
- 2 If another agent ate all food, then goal terminates with failure – agent is still hungry.
- 3 If agent ate some food (and thus completed *!eatFood*, but agent is still hungry, then more food will be eaten (*!eatFood* is re-tried).
- 4 If agent fails to eat pizza, then it should try eating something else! (*!eatFood* is re-tried)..

# Goal-programs Semantics

$$\frac{B \not\models \phi_s \vee \phi_f \quad \langle B, A, !e \rangle \longrightarrow \langle B', A', P \rangle}{\langle B, A, \text{Goal}(\phi_s, !e, \phi_f) \rangle \longrightarrow \langle B', A', \text{Goal}(\phi_s, P \triangleright P, \phi_f) \rangle} G_{\text{Adopt}}$$

$$\frac{B \not\models \phi_s \vee \phi_f \quad \langle B, A, P_1 \rangle \longrightarrow \langle B', A', P' \rangle}{\langle B, A, \text{Goal}(\phi_s, P_1 \triangleright P_2, \phi_f) \rangle \longrightarrow \langle B', A', \text{Goal}(\phi_s, P' \triangleright P_2, \phi_f) \rangle} G_{\text{Step}}$$

$$\frac{B \models \phi_s}{\langle B, A, \text{Goal}(\phi_s, P, \phi_f) \rangle \longrightarrow \langle B, A, \text{nil} \rangle} G_s \quad \frac{B \models \phi_f}{\langle B, A, \text{Goal}(\phi_s, P, \phi_f) \rangle \longrightarrow \langle B, A, ?\text{false} \rangle} G_f$$

$$\frac{P_1 \neq P_2 \quad B \not\models \phi_s \vee \phi_f \quad \langle B, A, P_1 \rangle \not\longrightarrow}{\langle B, A, \text{Goal}(\phi_s, P_1 \triangleright P_2, \phi_f) \rangle \longrightarrow \langle B, A, \text{Goal}(\phi_s, P_2 \triangleright P_2, \phi_f) \rangle} G_{\text{Restart}}$$

**Goal is adopted** –  $P = (\psi_1 : P_1, \dots, \psi_n : P_n)$  encodes all plans for  $e$ .

**Plan is executed** – only current strategy is updated!

**If  $\phi_s$  holds**, whole program terminates successfully

**If  $\phi_f$  holds**, whole program terminates with failure.

**If plan is finished or blocked**, original options are re-instantiated.

## Properties of $\text{Goal}(\phi_s, !e, \phi_f)$

- 1 Can be used by the BDI programmer to specify both **declarative** & **procedural** aspects of goals.
  - ▶ The whole original BDI reactive-online approach is maintained!;
- 2 Agent **never adopts** a goal that is already achieved or deemed impossible.
- 3 Agent **never drops** an adopted goal unless achieved or impossible.
  - ▶ Agent keeps trying and trying...
- 4 Agent **never pursues** goals which are achieved or impossible.
  - ▶ Will be terminated successfully or with failure.
- 5 Agent **never pursues** a goal that serves as a sub-goal for some higher-level motivating goal.
  - ▶ All sub-goals are removed when motivating goal is removed.

## Generating Top-Level Goals

How can the agent **self-generate** new top-level goals?

Equip our agents with a *motivation library*  $\mathcal{M}$  to accommodate the consideration of new goals in a proactive manner—the intrinsic agent's motivations or desires. Library  $\mathcal{M}$  consists of rules of the form:

$$\psi \rightsquigarrow \text{Goal}(\phi_s, !e, \phi_f),$$

If the agent comes to believe  $\psi$ , she should *consider adopting* the declarative goal-program  $\text{Goal}(\phi_s, !e, \phi_f)$ .

### Example

$\text{RoomDirty} \wedge \neg \text{Busy} \rightsquigarrow \text{Goal}(\neg \text{RoomDirty}, !\text{clean}, \text{HasWork}).$

## Generating Top-Level Goals (cont.)

To accommodate the motivational library we need to:

- 1 Extend agent-configurations to tuples of the form  $\langle \Pi, \mathcal{M}, \mathcal{B}, \mathcal{A}, \Gamma \rangle$
- 2 Add a new agent-level semantic rule:

$$\frac{\psi \rightsquigarrow \text{Goal}(\phi_s, !e, \phi_f) \in \mathcal{M} \quad \mathcal{B} \models \psi \quad \text{Goal}(\phi_s, P, \phi_f) \notin \Gamma}{\langle \Pi, \mathcal{M}, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Rightarrow \langle \Pi, \mathcal{M}, \mathcal{B}, \mathcal{A}, \Gamma \cup \{\text{Goal}(\phi_s, !e, \phi_f)\} \rangle} \text{Agt}_{\text{motiv}}$$

Observe that:

- ▶ Defined at the agent-level as it modifies the intention base.
- ▶ A completely new intention is created – new focus of attention.
- ▶ This mechanism has also been called elsewhere **agent desires** (3APL) or **automatic events** (JACK).

## Motivation

BDI-style **agent-oriented programming** is a novel approach to programming complex large systems in highly dynamic environments.

These systems are very open to changes in the environment:

- ▶ reactive – observe and respond to the environment;
- ▶ flexible – multiple options to achieve goals;
- ▶ robust – retry alternatives upon failure; commitment.

However, they usually lack any **lookahead planning** capabilities.

- ▶ Hypothetical reasoning about the future.

**HTN planning** is a well-known practical approach to planning.

Incorporate HTN planning into current BDI programming frameworks with:  
 (i) a clear semantics; (ii) a direct implementation.

# Basic Architecture of CAN


## Core Aspects of BDI Programming

**Beliefs:** information about the world.

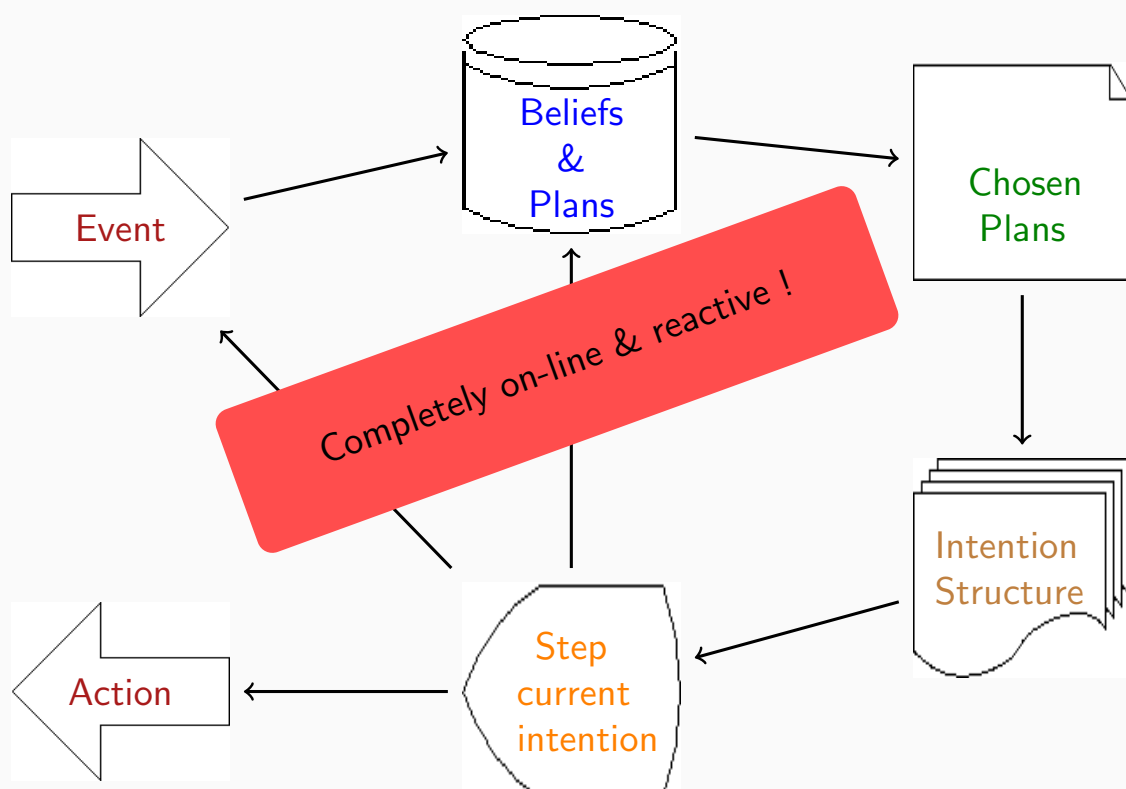
**Events:** goals/ desires to resolve; internal or external.

**Plan library:** recipes for handling goals-events.

**Intentions:** partially uninstantiated programs with commitment.



# The BDI Execution Cycle [Rao&Georgeff 92]



# Key Points of BDI Programming

BDI Progr. = Implicit Goal-based Programming + Rat. Online Executor

- ▶ Flexible and responsible to the environment: “reactive planning.” ✓
- ▶ Well suited for soft real-time reasoning and control. ✓
- ▶ Relies on context sensitive subgoal expansion: “act as you go.” ✓
- ▶ Leave for as late as possible the choice of which plans to commit to as the chosen course of action to achieve (sub)goals. ✓
- ▶ Modular and incremental programming. ✓
- ▶ Nondeterminism on choosing plans and bindings. ✓
- ▶ **BUT: No mechanism for doing lookahead for solving choices!** ⊗
  - ▶ Generally programmed *explicitly* by the BDI programmer.

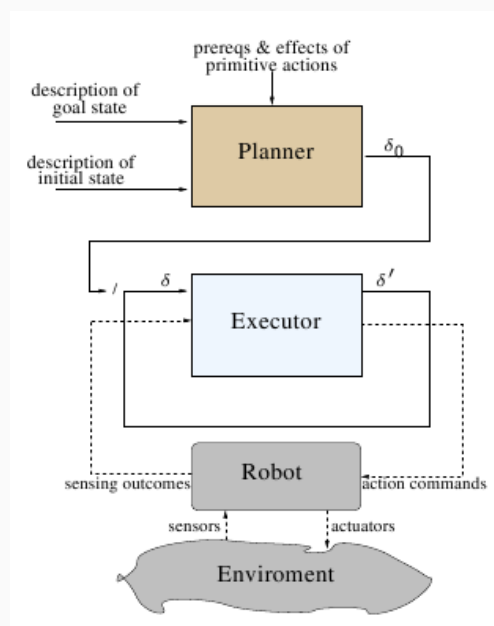
# Why do we want planning?

Some reasons for lookahead capabilities for BDI agents:

- ▶ Important resources may be used in taking actions that do not lead to successful outcome (e.g., fuel).
- ▶ Actions may sometimes not be reversible (e.g., break glass).
- ▶ Execution of actions may take substantial more time than just “thinking” ahead (e.g., traveling to a distant location).
- ▶ Actions’ execution may include undesirable side-effects (e.g., email notifications).

# Classical Planning [SRI Shakey "The Robot" '70]

- ▶ Planning from **first-principles**.
- ▶ Idea: achieve a goal by performing actions.
- ▶ **Goals-to-be**: *achieve  $\phi$  ( $At(airport)$ )*.
- ▶ Input to planner:
  - ▶ action descriptions;
  - ▶ initial state;
  - ▶ goal state.
- ▶ Output of planner:
  - ▶ a sequence of actions that achieves the goal.

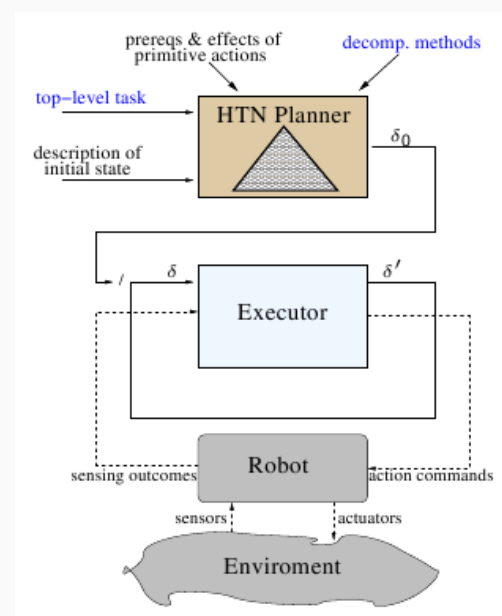


Obs.: planner has no domain information on how the goal state may be reached (e.g., common ways to go to the airport).

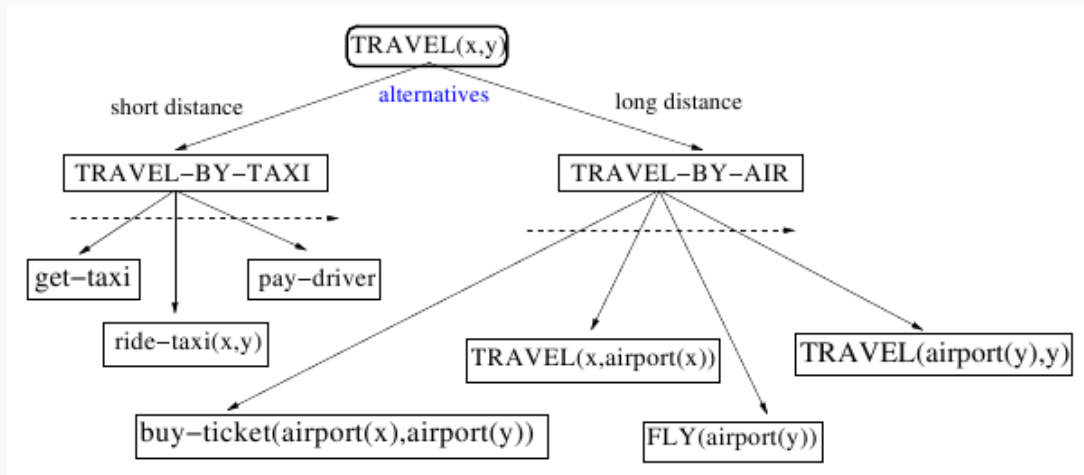
- ▶ planning algorithms don't take advantage of **all** problem structure.

# Hierarchical Task Networks (HTN) Planning

- ▶ Idea: accomplish some set of tasks, rather than to achieve a goal.
- ▶ **Goals-to-do**: *travelTo(*loc*)*.
- ▶ Includes domain procedural knowledge.
- ▶ Input to planner:
  - ▶ actions/operator descriptions;
  - ▶ initial state;
  - ▶ **top-level task**;
  - ▶ **decomposition methods**: how to decompose a task into a set of sub-tasks.
- ▶ Output of planner:
  - ▶ a sequence of actions that achieves the goal.
- ▶ Planning:
  - ▶ **Decompose tasks** by applying methods until all tasks in network are primitive actions.



# HTN Planning: An Example



**State:** set of atoms:  $At(loc)$ .

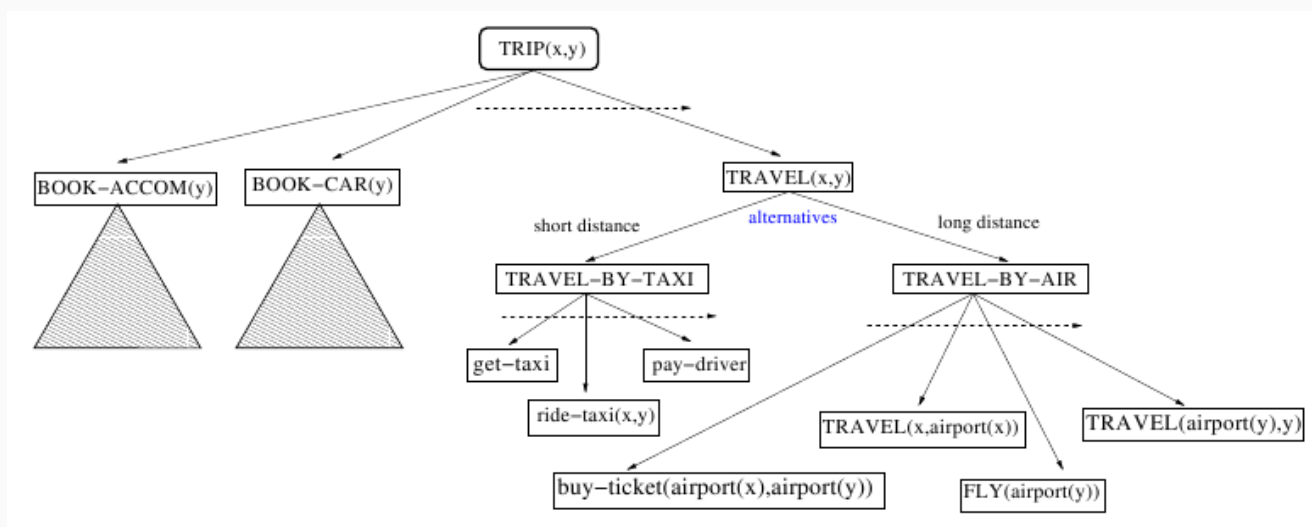
**Tasks:** **primitive** or **compound**.

**Task Network:** set of tasks  $T$  + order/state constraints  $\phi$ .

**Method:** a way to solve a compound task  $e$  using a network  $d$ .

**Plan:** a sequence of primitive tasks.

# HTN Planning: An Example (cont.)




## Key Points of HTN Planning

- ▶ **goals-to-do** instead of goals-to-be:  $goToLocation(YYZ)$  vs  $At(YYZ)$ .
- ▶ **Decomposition** of high-level **tasks**.
- ▶ User provides (procedural) knowledge capturing lots of **useful domain information**:
  - ▶ Methods correspond to “**recipes** for accomplishing things.”;
  - ▶ **If we already have an idea of how to solve a problem, we might as well use it in our planner!**
- ▶ Well understood **semantics** and **implementations**.
  - ▶ Operational and model-theoretic semantics.
  - ▶ Implementations: SHOP, SHOP2, UMCP, JSHOP, etc.
- ▶ HTN planners have been some of the most successful & commonly used.
- ▶ Some similarities with “programming”.
- ▶ Subsumes first-principle “STRIP” planning.

## Components of HTN Planning

- ▶ **State**: set of atoms + CWA
- ▶ **Task**: a common goal that the agent may need to achieve.
  - ▶ **Primitive task  $a$** : with precondition and effects:  $ride(x, y)$ ,  $payTaxi$ , etc.
  - ▶ **High-level task  $e$** : cannot be directly executed:  $travel(x, y)$
- ▶ **Task Network  $d = (T, \phi)$** : set of tasks  $T$  + order/state constraints  $\phi$ .  
 $d_{travel} = ( \{ t_4 = travel(a(y), y), t_2 = travel(x, a(x)), t_3 = fly(a(x), a(y)), t_1 = getTicket(a(x), a(y)) \}, (t_1 \prec t_3 \wedge t_2 \prec t_3 \wedge t_3 \prec t_4) )$



- ▶ **Method  $m = [e, \psi, d]$** : “recipe” to solve task  $e$ .  
 $[travel(x, y), LongDistance(x, y), d_{travel}]$



## Formal Specification of HTN-Planning [Erol et al 94]

- ▶ An **HTN planning domain** is a pair  $\mathcal{D} = (\Pi, \Lambda)$ :
  - ▶ method library  $\Pi$ ;
  - ▶ STRIP-like (add and delete lists) action description  $\Lambda$ .
- ▶ An **HTN problem** is a triple  $\mathbf{P} = \langle d, \mathcal{B}, \mathcal{D} \rangle$ :
  - ▶ solve task network  $d$  in state  $\mathcal{B}$  using domain  $\mathcal{D}$ .
- ▶ A **plan**  $\sigma$  is a sequence of primitive tasks.
 

*getCab · ride(rome, a(rome)) · checkIn · walk(gate) ...*
- ▶ An **HTN planning solution** for problem  $\mathbf{P}$  is a plan  $\sigma$  that stands for a *full successful decomposition* of task network  $d$ .
  - ▶  $sol(d, \mathcal{B}, \mathcal{D})$ : set of all plans that solve  $d$ .

## The CAN Language [Winikioff et al. 2002]

A CAN agent is defined as  $\text{Agt} = \langle \mathcal{N}, \Pi, \mathcal{M}, \mathcal{B}, \mathcal{A}, \Gamma \rangle$ , where:

- ▶  $\mathcal{N}$  is the **agent name**.
- ▶  $\mathcal{B}$  is the **belief base**: current agent's knowledge.
- ▶  $\mathcal{A}$  is the sequence of actions executed so far.
- ▶  $\Pi$  is a **plan library** containing plan rules  $e : \psi \leftarrow P$ :
  - ▶  $e$  is the **triggering event**;  $\psi$  is the **context condition**;  $P$  is the **plan-body**:  
 $P ::= nil \mid act \mid ?\phi \mid +b \mid -b \mid P_1; P_2 \mid P_1 \parallel P_2 \mid !e \mid \underline{Goal(\phi_s, P_1, \phi_f)}$
- ▶  $\mathcal{M}$  is the motivation library.
- ▶  $\Gamma$  is the **intention base**: partially uninstantiated plan-bodies.

# The CAN Language

However, CAN does not include any planning mechanism.

**The objective:** add lookahead capabilities to CAN in a way that:

- 1 the original semantics can be extended;
- 2 an implementation can be easily produced.

We will extend CAN with:

- 1 A library of STRIP-like action descriptions  $\Lambda$  (precondition & effects);

$$a : \psi \leftarrow \Phi^-; \Phi^+ \in \Lambda$$

$$pickUp(x) : Clear(x) \leftarrow \{Clear(x)\}^-; \{Holding(x)\}^+ \in \Lambda$$

- 2 A new programming construct **Plan(*P*)**.

► Plan(*P*): look/check for a full solution for BDI plan *P*.

observe similarity with  $\Sigma$  in IndiGolog!

# BDI Systems and HTN Planners: Similarities

BDI Systems	HTN Systems
belief base	state
plan library	method library
event	high-level task
plan rule	method
plan-body/program	network task
plan rule context	method precondition
action	primitive task
test ? <i>p</i> in plan-body	state constraints
sequence in plan-body	ordering constraint $\prec$
parallelism in plan-body	no ordering constraint

$[travel(x, y), LongDist(x, y), d_{travel}]$

$d_{travel} = ( \{ t_4 = travel(a(y), y), t_2 = travel(x, a(x)), t_3 = fly(a(x), a(y)), t_1 = getTicket(a(x), a(y)) \}, (t_1 \prec t_3 \wedge t_2 \prec t_3 \wedge t_3 \prec t_4) )$

$travel(x, y) : LongDist(x, y) \leftarrow (!getTicket(a(x), a(y)) \parallel !travel(x, a(x))); !fly(a(x), a(y)); !travel(a(y), y)$

# The CANPlan Language: Operational Semantics

Recall semantics of CAN is modularly defined in two levels:

- 1 Agent-level semantics:  $\langle \mathcal{N}, \Pi, \mathcal{M}, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{M}, \mathcal{B}', \mathcal{A}', \Gamma' \rangle$ .
- 2 Intention-level semantics:  $\langle \Pi, \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \langle \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$ .

We now use 2 labelled intention-level transitions:

- ▶  $\xrightarrow{\text{bdi}}$ : transitions for the BDI execution cycle – all rules as before.
- ▶  $\xrightarrow{\text{plan}}$ : transitions when planning – **almost** all the rules as before.

We also add the STRIP-like action library  $\Lambda$ :

- 1 Agent-level semantics:  
 $\langle \mathcal{N}, \Lambda, \Pi, \mathcal{M}, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Lambda, \Pi, \mathcal{M}, \mathcal{B}', \mathcal{A}', \Gamma' \rangle$ .
- 2 Intention-level semantics:  
 $\langle \Lambda, \Pi, \mathcal{B}, \mathcal{A}, P \rangle \xrightarrow{\text{bdi}} \langle \Lambda, \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$  and  
 $\langle \Lambda, \Pi, \mathcal{B}, \mathcal{A}, P \rangle \xrightarrow{\text{plan}} \langle \Lambda, \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$ .

# The CANPlan Language: Operational Semantics

New construct  $\text{Plan}(P)$ : perform off-line lookahead on plan  $P$ .

$$\frac{\langle \mathcal{B}, \mathcal{A}, P \rangle \xrightarrow{\text{plan}} \langle \mathcal{B}', \mathcal{A}', P' \rangle \quad \langle \mathcal{B}', \mathcal{A}', P' \rangle \xrightarrow{\text{plan}_*} \langle \mathcal{B}'', \mathcal{A}'', \text{nil} \rangle}{\langle \mathcal{B}, \mathcal{A}, \text{Plan}(P) \rangle \xrightarrow{\text{bdi}} \langle \mathcal{B}', \mathcal{A}', \text{Plan}(P') \rangle} \text{Plan}$$

Inspired by the semantics of  $\Sigma(\delta)$  in IndiGolog  
(De Giacomo & Levesque 99)

# The CANPlan Language: Operational Semantics (cont.)

We also need to do the following three changes:

- 1 account for actions' preconditions and effects;

$$\frac{a : \psi \leftarrow \Phi^-; \Phi^+ \in \Lambda \quad a\theta = act \quad \mathcal{B} \models \psi\theta}{\langle \mathcal{B}, \mathcal{A}, act \rangle \longrightarrow \langle (\mathcal{B} \setminus \Phi^-) \cup \Phi^+, \mathcal{A} \cdot act, nil \rangle} do$$

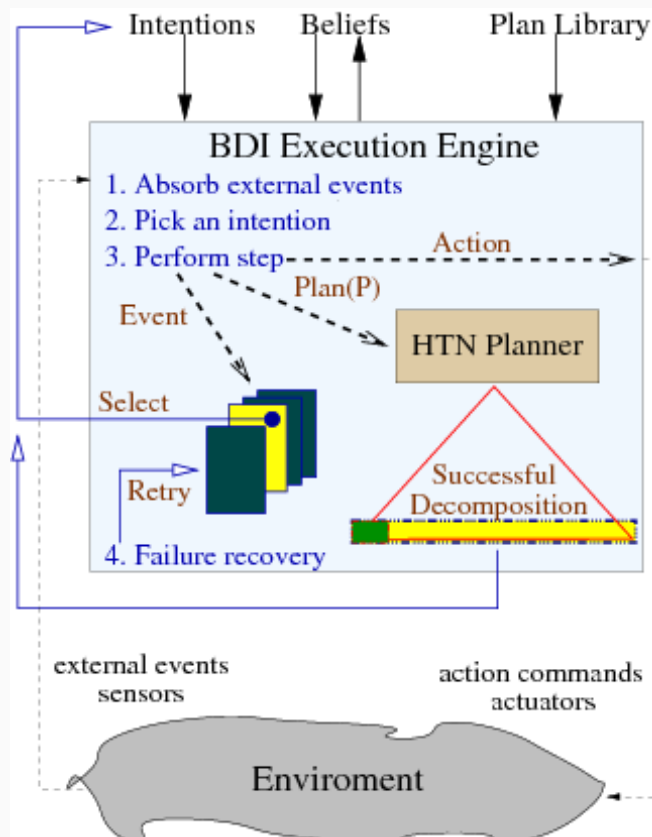
- 2 allow retry of alternatives *only* in the BDI execution cycle;

$$\frac{P_1 \neq nil \quad \langle \mathcal{B}, \mathcal{A}, P_1 \rangle \not\rightarrow \langle \mathcal{B}, \mathcal{A}, P_2 \rangle \longrightarrow \xrightarrow{bdi} \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \xrightarrow{bdi} \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle} \triangleright_f$$

- 3 perform agent steps based on bdi-type basic transitions.

$$\frac{P \in \Gamma \quad \langle \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \xrightarrow{bdi} \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', (\Gamma \setminus \{P\}) \cup \{P'\} \rangle} Agt_{step}$$

## CANPlan Architecture



## Some Remarks About Plan

- 1 Plan is **built-in** within the BDI execution cycle.
  - ▶ on demand look-ahead planning.
- 2 Plan uses exactly the **same domain knowledge** as given to the BDI system.
  - ▶ no separate knowledge needs to be given!
  - ▶ because BDI and HTN system work on the same kind of information...
- 3 Plan's semantics is defined **"in terms" of the BDI cycle**.
  - ▶ defined in terms of the same semantic rules;
  - ▶ because BDI and HTN system work in a similar way...
- 4 But **failure-handling** is left to the BDI cycle only!

## Main Results About Plan( $P$ ) (cont.)

### Theorem

*An intention Plan( $P$ ) **never fails** in an environment-free execution.*

### Theorem

*All executions of Plan( $P$ ) correspond to **HTN solutions** of  $P$  w.r.t. the agent's plan-library  $\Pi$  (and vice-versa).*

Basis for an implementation: JACK+ JSHOP (Lavindra & Padgham 05):

- ▶ JACK is a BDI-style programming language similar to CAN;
- ▶ JSHOP is an Java-based HTN planner implementation.

(no concurrency || and no goal-programs Goal, yet)

## Planning for Declarative Goals: Use $P$ for achieving $\phi_s$

- (A)  $P$  is used towards the eventual satisfaction of goal  $\phi_s$ .
- (B) Allow partial executions of  $P$  achieving  $\phi_s$ .
- (C) Commitment on the goal  $\phi_s$ .
- (D) Goal dropping mechanism.
- (E)  $P$  may be solved partially.

ALTERNATIVES	A	B	C	D	E
$\text{Plan}(P; ?\phi_s)$	✓				
$\text{Plan}(\text{Goal}(\phi_s, P, \phi_f))$	✓	✓		✓	
$\text{Goal}(\phi_s, \text{Plan}(P), \phi_f)$		✓	✓	✓	
$\text{Goal}(\phi_s, \text{Plan}(P; ?\phi_s), \phi_f)$	✓	✓	✓	✓	
$\text{Goal}(\phi_s, \text{Plan}(\text{Goal}(\phi_s, P, \phi_f)), \phi_f)$	✓	✓	✓	✓	✓

$$\text{Plan}(\phi_s, P, \phi_f) \stackrel{\text{def}}{=} \text{Goal}(\phi_s, \text{Plan}(\text{Goal}(\phi_s, P, \phi_f)), \phi_f)$$

## Other Open Issues

- 1 Plan **monitoring & replanning**:
  - ▶ What to do when a planning solution fails?
- 2 Planning in the context of **other intentions**.
  - ▶ Reasoning about negative & positive interactions.
- 3 Planning up to some level of **abstraction**.
  - ▶ Leave details to the BDI execution engine.
- 4 Using **first-principles planning** for learning new plans:
  - ▶ When no relevant plan is available  $\implies$  classical planning.
  - ▶ Add new plan to the BDI plan library  $\Pi$ .
- 5 ...

## Review

In this lecture we have shown how to

**1 Accommodate declarative goals in CAN:**

- ▶ motivated the need for declarative goals (e.g., decoupling goal achievement vs plan termination).
- ▶ incorporated new construct  $\text{Goal}(\phi_s, P, \phi_f)$  into the language.

**2 Accommodate on-demand planning in CAN:**

- ▶ motivated the need for planning in BDI systems;
- ▶ reviewed HTN-planning: off-line decomposition of tasks;
- ▶ incorporated new construct  $\text{Plan}(P)$  to the language;






## Next Lecture

# JACK Agent-Oriented Programming Language

## Declarative Goals

- 
 Michael Winikoff, Lin Padgham, James Harland, and John Thangarajah.  
 Declarative & procedural goals in intelligent agent systems.  
*In Proceedings of the KR-02, 2002.*
- 
 Birna van Riemsdijk, Mehdi Dastani, Frank Dignum and John-Jules Ch. Meyer  
 Dynamics of Declarative Goals in Agent Programming.  
*In Proceedings of the DALT-05, 2005.*
- 
 Jomi Hubner, Rafael, Bordini and Michael Wooldridge.  
 Programming declarative goals using plan patterns.  
*In Proceedings of the DALT-06, 2006.*
- 
 Frank S. de Boer, Koen V. Hindriks, Wiebe van der Hoek, and John-Jules Ch. Meyer.  
 A Verification Framework for Agent Programming with Dec. Goals.  
*Journal of Applied Logic, 5(2):277–302, 2007.*

## Planning in BDI Systems

- 
 Kutluhan Erol, James Hendler, and Dana S. Nau.  
 Complexity Results for HTN Planning  
*Annals of Math in Artificial Intelligence, 18(1):69-93, 1996*
- 
 D. Nau, Y. Cao, A. Lotem, and H. Muñoz Avila.  
 Shop: Simple hierarchical ordered planner.  
*In Proceedings of IJCAI-99, pages 968–973, 1999.*
- 
 Olivier Despouys and Francois Felix Ingrand.  
 Propice-plan: Toward a unified framework for planning and execution.  
*In Proc. of European Conference on Planning, pages 278–293, 1999.*
- 
 Jürgen Dix, Héctor Muñoz-Avila, Dana S. Nau, and Lingling Zhang.  
 IMPACTing SHOP: Putting an AI planner into a multi-agent env.  
*Annals of Mathematics and Artificial Intelligence, 37(4):381–407, 2003.*
- 
 Sebastian Sardina, Lavindra P de Silva, and Lin Padgham.  
 Hierarchical planning in BDI agent programming languages: A formal approach.  
*In Proceedings of AAMAS-06, pages 1001–1008, 2006.*