

# Lecture 7: Formal Models of BDI Programming

Autonomous Agents and Multiagent Systems  
DIS, La Sapienza - PhD Course

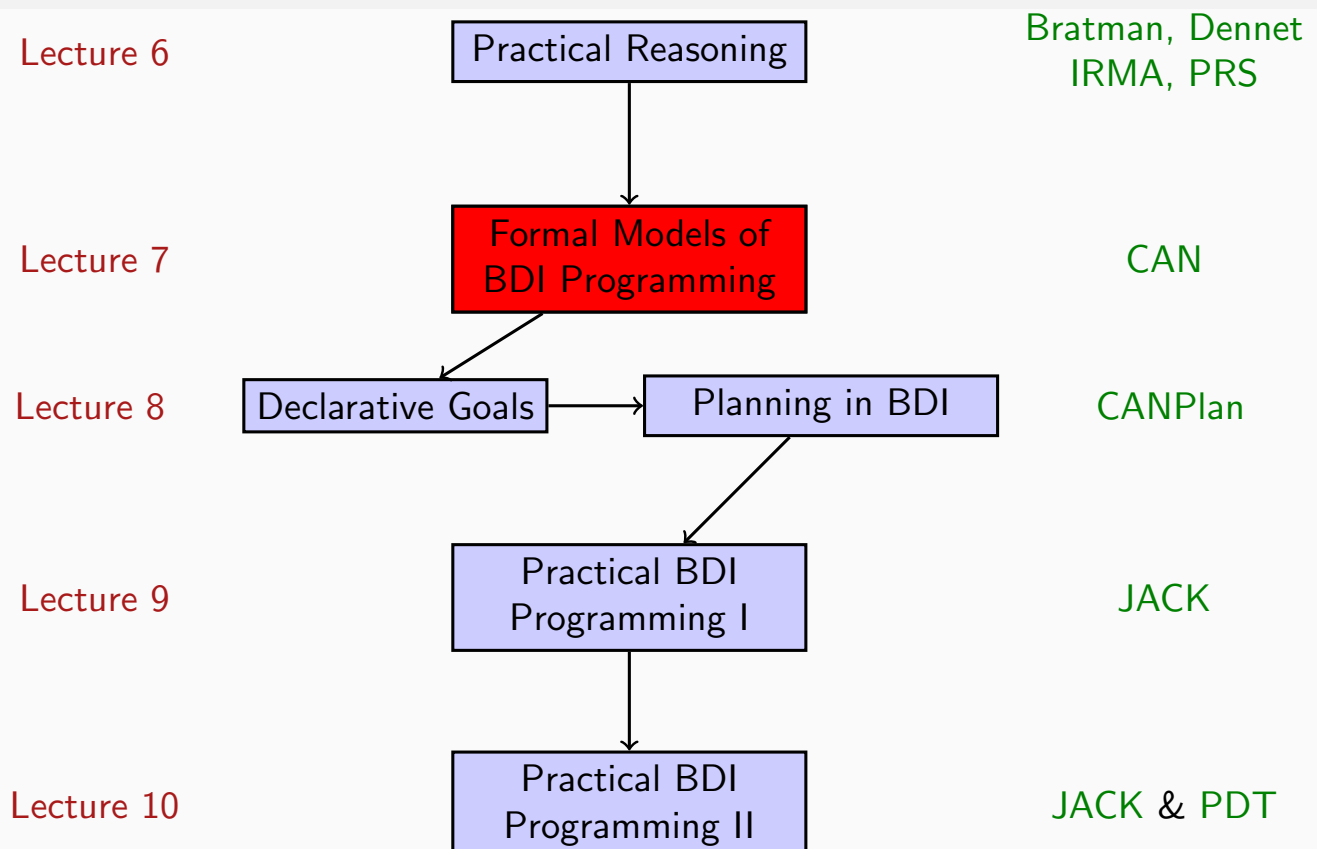
Sebastian Sardina<sup>1</sup>

<sup>1</sup>Department of Computer Science and Information Technology  
RMIT University  
Melbourne, AUSTRALIA



November 27, 2007

## Roadmap for Next Lectures



# Outline

- 1 Introduction
- 2 BDI Framework
  - Core Aspects
  - Overview
- 3 The CAN Language
  - Overview
  - Syntax
  - Semantics
  - Other Languages
- 4 Conclusions

## Review of Last Lecture: Practical Reasoning

- 1 **Intentional stance** (Daniel Dennett (1987))
  - ▶ High-level **abstraction** of behavior at the level of **minds**.
  - ▶ Rational behavior can be understood in terms of **mental properties**:
    - ▶ beliefs, desires, goals;
    - ▶ fear, hopes, etc.
- 2 **Practical reasoning** (Michael Bratman (1990))
  - ▶ Reasoning for acting: the process of figuring out what to do.
  - ▶ Two activities: **deliberation** and **means-end analysis**.
- 3 **Commitments** (on goals/intentions & plans)
  - ▶ fanatical, single-minded, open-minded.
- 4 **Agent architectures**
  - ▶ IRMA & PRS.
  - ▶ Built around: beliefs, desires, plan libraries, intentions, filter, etc.
- 5 **Agent theory**
  - ▶ Cohen & Levesque's "Intention = Choice+Commitment".
  - ▶ Rao & Georgeff's BDI logic.

# This Lecture: BDI Programming

Objective: a **programming language** that can provide:

**autonomy**: does not require continuous external control;

**pro-activity**: pursues goals over time; goal directed behavior;

**situatedness**: observe & act in the environment;

**reactivity**: perceives the environment and responds to it.

**flexibility**: achieve goals in several ways.


**robustness**: will try hard to achieve goals.

And also: **modular scalability** & **adaptability**!

## An Example: Gold Mining Game

2 teams competing to collect and drop gold in the depot

- ▶ dynamic
- ▶ complex
- ▶ unknown information
- ▶ failing actions
- ▶ failing sensors
- ▶ multi-agents



## Some Constraints

We want to **program intelligent systems** under the following constraints:

- 1 The agent **interacts** with an external environment.
  - ▶ A grid world with gold pieces, obstacles, and other agents.
- 2 The environment is (highly) **dynamic**; may change in unexpected ways.
  - ▶ Gold pieces appear randomly.
- 3 Things can go wrong; plans and strategies may **fail**.
  - ▶ A path may end up being blocked.
- 4 Agents have **dynamic** and **multiple** objectives.
  - ▶ Explore, collect, be safe, communicate, etc.
  - ▶ Motivations/goals/desires may come and go.

## Some Assumptions

Luckily, we can also assume that:

- 1 **Failure** is generally not catastrophic.
  - ▶ If gold is dropped, we just pick it up again.
  - ▶ If a tree blocks the path, we just go around it.
- 2 We can understand the system at the “intentional” level.
  - ▶ Agents “desire” to collect as much gold as possible.
  - ▶ Agents “believe” they are close to the depot.
- 3 There is “some” sensible known **procedural knowledge** of the domain.
  - ▶ It is “good” to avoid obstacles.
  - ▶ If we see gold close, then go and collect it.
  - ▶ If we bump into an unknown wall, then walk along it.

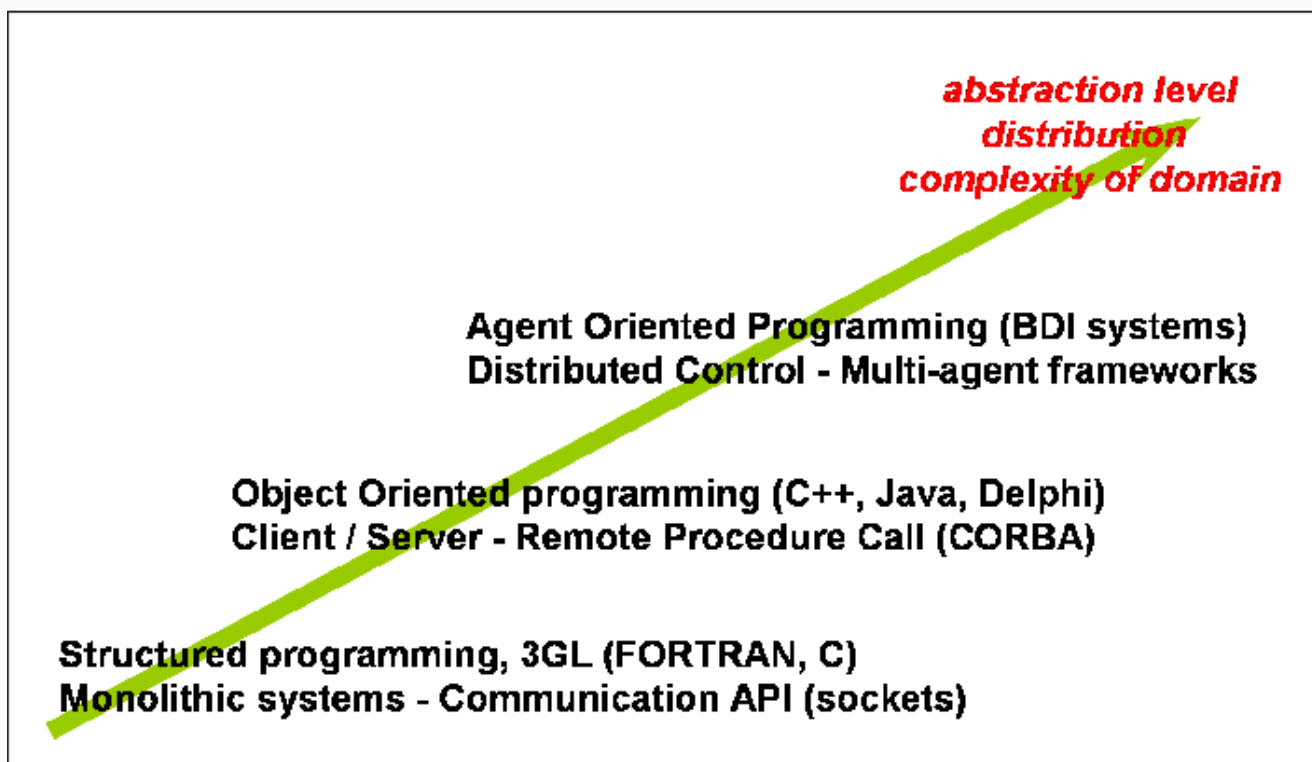
# Agent-Oriented Programming (Shoam 1993)

- 1 Philosophers have produced theories of (human) rational action:
  - ▶ Folk psychology.
  - ▶ Practical reasoning.
  - ▶ Intentional systems.
- 2 Theorists have taken this and developed theories to represent the properties of agents (humans or not).
  - ▶ Relation between mental attitudes.
  - ▶ Commitment.
  - ▶ Rational architectures.

So, why not **directly program agents in terms of the mentalistic, intentional notions?**

∴ we will study one agent-oriented approach: BDI-style Programming.

## Technology Development



## Parent Technologies

- 1 artificial intelligence
- 2 software engineering
- 3 distributed systems
- 4 organizational science
- 5 databases
- 6 economics
- 7 game theory
- 8 artificial life

## Some Agent Companies

- 1 Agent Oriented Software (JACK Intelligent Agents)
- 2 Reticular (AgentBuilder intelligent agents)
- 3 Gensym Corp. (G2: real-time business rule engine platform)
- 4 Agentis Software (AdaptivEnterprise: goal-oriented system)
- 5 Whitestein (autonomic self-managing, goal-oriented business solutions)
- 6 IBM (Aglets: Java based mobile agent platform)
- 7 Genesys Telecommunications (customer service)
- 8 Hewlett Packard
- 9 ...

## Core Aspects

Most basic concepts:

**Agent** Autonomous SW entity.

**Percepts** Information **perceived** from the environment.

**Actions** **Affects** the environment.

Some additional concepts:

- 1 Message (inter agent communication)
- 2 Event (trigger)
- 3 Goal (what to do)
- 4 Plan (how to do)
- 5 Interaction protocol (conversation pattern)
- 6 Organization (also team, institution)
- 7 Role (agent abstraction)
- 8 ...

## Building Agents: BDI Agent-oriented Programming

A new **programming model** and **architecture** to simplify the construction of todays **large complex systems situated in dynamic environments**:

- 1 View a system as composed of autonomous interacting entities (**agents**) which pursue their **own goals** and act in a **rational** manner.
- 2 Internal state and decision process of agents is modelled in an intuitive manner following the notion of **mental attitudes**.
- 3 **Goal orientation**: instead of directly requesting the agents to perform certain actions, the developer can define more **abstract goals** for the agents.
  - provides a certain degree of flexibility on how to achieve the goals.

Can be seen as a “successor” of object-oriented programming

# Some BDI Agent-oriented Programming Languages

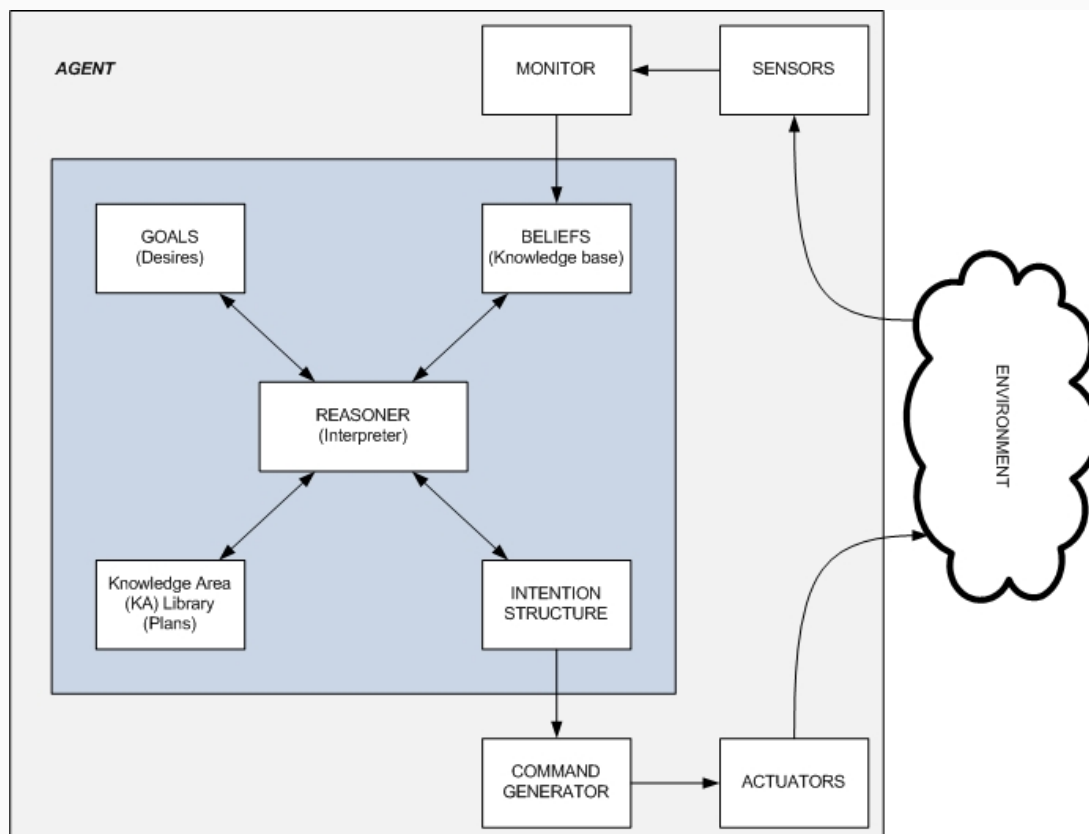
Some **formal** BDI programming languages:

- 1 AgentSpeak: first formal BDI language.
- 2 3APL: language with different kind of plan rules.
- 3 GOAL: based on declarative goals only.
- 4 **CAN(Plan): failure-handling, declarative goals, & planning.**  $\Leftarrow$  **TODAY**

Some BDI programming language **systems/platforms/architectures**:

- 1 PRS and dMars: first BDI-based systems; C++ based.
- 2 JAM: a hybrid extension of PRS.
- 3 JASON: JAVA-based implementation of AgentSpeak.
- 4 JADEX: based on JADE communication platform.
- 5 SPARK: SRI's BDI system.
- 6 **JACK: powerful commercial JAVA-based BDI-system.**  $\Leftarrow$  **NEXT LECT.**

# Typical BDI-style System





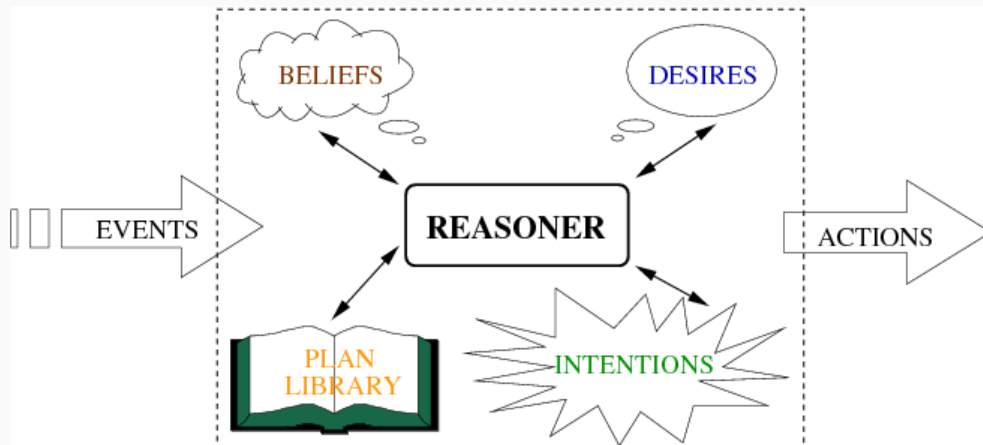
# Key Features of BDI Agent-oriented Systems

**Beliefs:** information about the world.

**Events:** goals/desires to resolve; internal or external.

**Plan library:** recipes for handling goals-events.

**Intentions:** partially uninstantiated programs with commitment.



# Key Features of BDI Agent-oriented Systems (cont.)


In the gold-mining game:

**Beliefs:** current location & location of depot.  
size of grid, # of gold pieces carrying, etc.

**Events:** a gold piece is observed east;  
player3 communicates its location;  
the coordinator requests to explore the grid;  
we formed the *internal goal* to travel to *loc(10, 22)*

**Plan library:** if I see gold here & I am not full, collect it.  
if I hit an obstacle, go around it.  
if I don't know of any gold, explore grid.  
if I see gold around, move there and collect.

**Intentions:** I am currently traveling to the depot.  
I am informing my team-mates of new obstacles I find.



# Events & Plans

## 1 Events stand for the goals/desires/tasks to be achieved or resolved:

- ▶ **percepts:** *goldAt(east)*, *goldDropped*, etc;
- ▶ **communication:** *told(player3, loc(3, 2))*;
- ▶ **external request/goal:** *achieve(explore\_grid)*;
- ▶ **internal sub-goal:** *go\_to(loc(10, 22))*.

## 2 Plans stand for strategies useful to resolve (pending) events:

- ▶ encode typical operational procedures in the domain;
- ▶ non-deterministic;
- ▶ event & context dependent;

$$e : \psi \longleftarrow P$$

$P$  is a good strategy to resolve event  $e$  if context  $\psi$  is believed true.

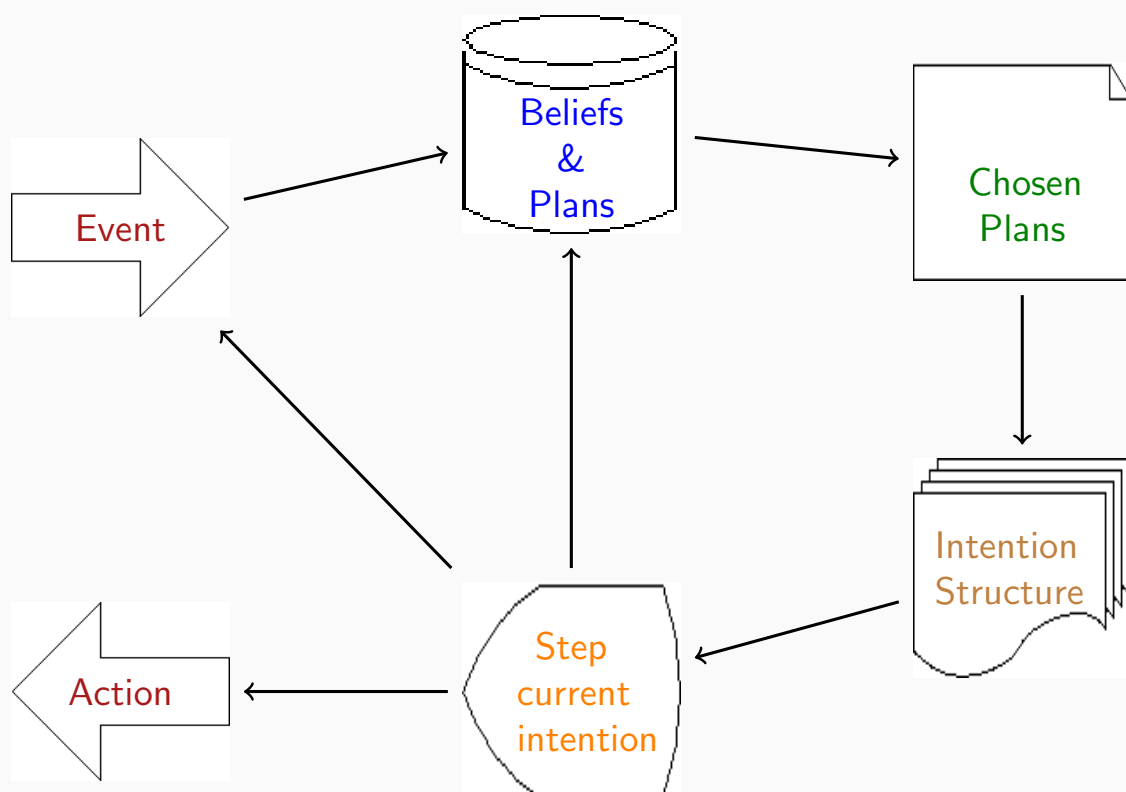
# Plans in PRS: Clearing a Block

```
Plan: {
NAME: "Clear a block"
GOAL:
    ACHIEVE CLEAR $OBJ;
CONTEXT:
    FACT ON $OBJ2 $OBJ;
BODY:
    EXECUTE print "Clearing " $OBJ2 " from on top of " $OBJ "\n";
    EXECUTE print "Moving " $OBJ2 " to table.\n";
    ACHIEVE ON $OBJ2 "Table";
EFFECTS:
    EXECUTE print "CLEAR: Retracting ON " $OBJ2 " " $OBJ "\n";
    RETRACT ON $OBJ1 $OBJ;
FAILURE:
    EXECUTE print "\n\nClearing block " $OBJ " failed!\n\n";
}
```

# Intentions

- 1 Agent's intentions are determined **dynamically** by the agent at **runtime** based on its known facts, current goals, and available plans.
- 2 An intention is just a **partially executed** strategy:
  - ▶ comes from the plan library when resolving events.
- 3 An intention represent a **focus of attention**:
  - ▶ something the agent is currently working on;
  - ▶ actions/behavior arises as a consequence of executing intentions.
- 4 An agent may have **several intentions** active at one time.
  - ▶ different simultaneous focuses of attention;
- 5 A **new intention is created** when an external event is addressed.
- 6 An intention may create/post an **internal event**:
  - ▶ the intention will be updated when this event is addressed.

## The BDI Execution Cycle [Rao&Georgeff 92]



# The BDI Execution Cycle [Rao&Georgeff 92]



## The BDI Execution Cycle: Detailed Version

- 1 **Observe** the environment for new *external* events.
- 2 Pick a pending **event** *e*.
- 3 Select **relevant** plans from library (match event).
- 4 Select **applicable** plans from relevant set (match context).
- 5 If event *e* is **external**, create new intention with selected plan.
- 6 If event *e* is **internal**, update intention with selected plan on top.
- 7 Partially **execute** some intention (may post internal events).
  - ▶ If execution fails, then perform **failure recovery**.
- 8 Repeat cycle.

# Key Points of BDI Programming

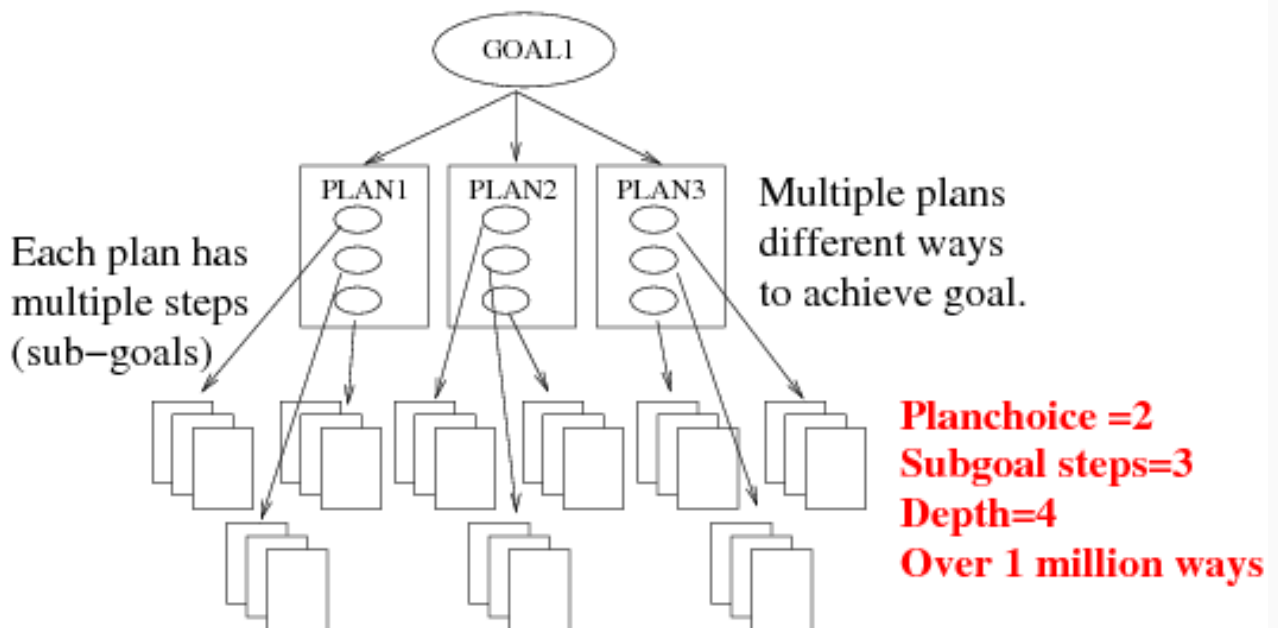
- ▶ Flexible and responsible to the environment: “**reactive planning**.”
  - ▶ Well suited for **soft real-time** reasoning and control.
- ▶ Relies on **context sensitive** subgoal expansion: “act as you go.”
- ▶ Leave for **as late as possible** the choice of which plans to commit to as the chosen course of action to achieve (sub)goals.
- ▶ **Modular** and **incremental** programming.
- ▶ Nondeterminism on **choosing** plans and bindings.

# Key Points of BDI Programming

- ▶ Flexible and responsible to the environment: “reactive planning.”
  - ▶ Well suited for soft real-time reasoning and control.
- ▶ Relies on context sensitive subgoal expansion: “act as you go.”
- ▶ Leaves the user in control of the reasoning process.
- ▶ Modular and incremental programming.
- ▶ Nondeterminism on choosing plans and bindings.

BDI Programming =  
Implicit Goal-based Programming + Rational Online Executor

# Possibility of Many Options

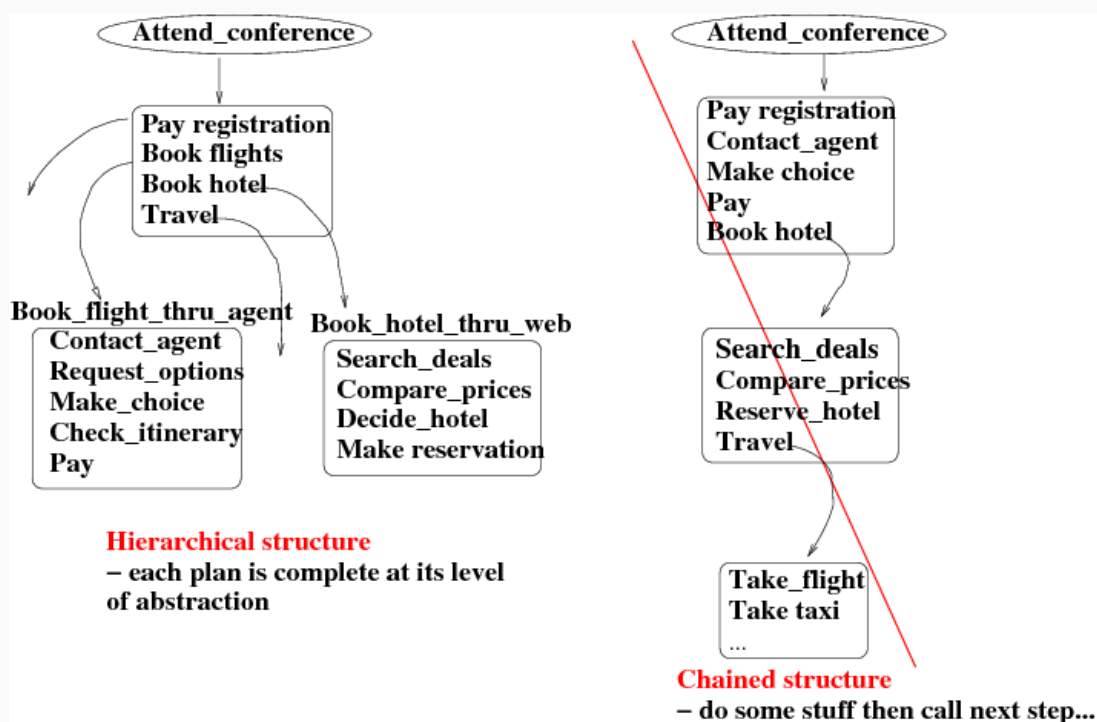


Here we have 30 plans, 81 way to achieve the goal.  
depends on choice of plans, number of steps, and  
depth of tree:

# Making Use of the BDI Framework

- 1 Provide alternative plans where possible.
- 2 Break things down into subgoal steps.
- 3 Use subgoals and alternative plans rather than **if... then** in code.
- 4 Keep plans small and modular.
- 5 Plans are abstract modules - don't chain them together like a flowchart.

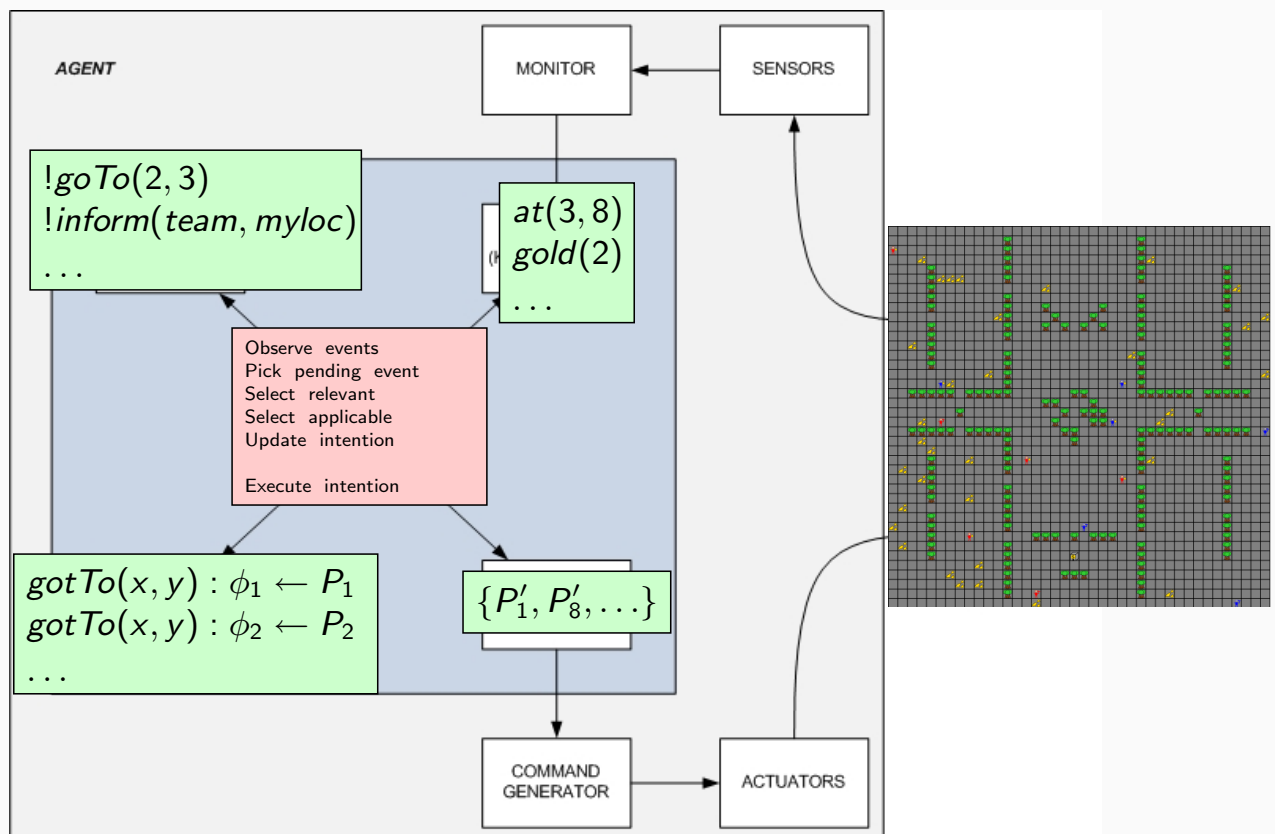
## Plan Structure



# Structuring Plans and Goals

- 1 Make each plan **complete** at a particular **abstraction level**.
  - ▶ A high-level but complete plan for *Attend\_Conference*.
- 2 Use a **subgoal** - even if only one plan choice for now.
  - ▶ Decouple a goal from its plans.
- 3 Modular and easy to **add other plan choices** later.
  - ▶ Booking a flight can now be done with the Internet, if available!
- 4 Think in terms of **subgoals**, not function calls.
  - ▶ What way-points do we need to achieve so as to realize a goal?
- 5 Learn to **pass information** between subgoals.
  - ▶ How are these way-points inter-related w.r.t. data?


# Typical BDI-style System



# Historical Evolution

- 1 A few actual BDI systems-frameworks were developed (late 80s'):
  - ▶ IRMA, PRS, dMars, etc.
- 2 Anand Rao tried to formally capture the common features (1996):
  - ▶ AgentSpeak appeared: events, beliefs, intentions, rational cycle, etc.
  - ▶ Can be seen as an elegant extension of logic programming for the implementation of BDI agents (reactive planning systems).
- 3 Many other languages appeared in the BDI-tradition (late 90s'):
  - ▶ 3APL
  - ▶ JACK
  - ▶ JADEX
  - ▶ JASON
  - ▶ ...
- 4 CAN(Plan) is similar to AgentSpeak but with (2002-2007):
  - ▶ built-in semantics for failure handling (as in real BDI systems);
  - ▶ declarative goals (as in agent theories);
  - ▶ planning capabilities (as in HTN-planning & IndiGolog).
  - ▶ interleaved concurrency in plans (as in ConGolog);

# Basic Architecture of CAN





# Basic Syntax of CAN

- ▶ Beliefs: *predicate(terms)* (e.g., *curLoc(2,3)*)
- ▶ Actions: *action(terms)* (e.g., *pickGold, move(left)*).
- ▶ Events: *event(terms)* (e.g., *goTo(10,20), block(agt3)*).
- ▶ Plans: *event : context*  $\leftarrow$  *body*.
- ▶ Body:

<i>act</i>	<i>primitive action</i>
<i>+b</i>	<i>belief addition</i>
<i>-b</i>	<i>belief deletion</i>
<i>?<math>\phi</math></i>	<i>tests goal</i>
<i>!e</i>	<i>posting of achievement event goal</i>
<i>P<sub>1</sub>; P<sub>2</sub></i>	<i>sequence</i>
<i>P<sub>1</sub>  P<sub>2</sub></i>	<i>interleaved concurrency</i>

## A Simple Example

```
friend(john).
friend(anne).
time(3pm).
phone(john,93812298).
....
```

```
meet(X) : friend(X) <- !greet_friend(X).
meet(X) : not friend(X) <- !greet(X).
```

```
greet(X) : time(T) and T<=12pm <- say('Good morning').
greet(X) : time(T) and T>12pm <- say('Good afternoon').
greet(X) : true <- say('Hello!').
```

```
greet_friend(X) : true <- say('Hi '+X).
```

```
told(X,phone,N) : not friend(X) <- +friend(X), +phone(X,N).
```

## A Simple Example II

```

currLoc(10,12).
goldAt(10,13).
goldAt(2,3).
carrying(2).
depotLoc(20,14).
...

goTo(X,Y) : grid_known >50 <- !planPath(X,Y, Path), !follow(Path).
goTo(X,Y) : not grid_known>50 <- !getTowards(X,Y,X2,Y2),
                                !moveTo(X2,Y2).

...

act : fullLoaded <- ?depotLoc(X,Y), +toDepot, !goTo(X,Y).
act : currLoc(X,Y) and goldAt(X,Y) <- pick.
act : goldAt(X,Y) <- !goTo(X,Y).
act : not goldAt(X,Y) <- ?team(Agnt), !askForGold(Agnt).
...

```

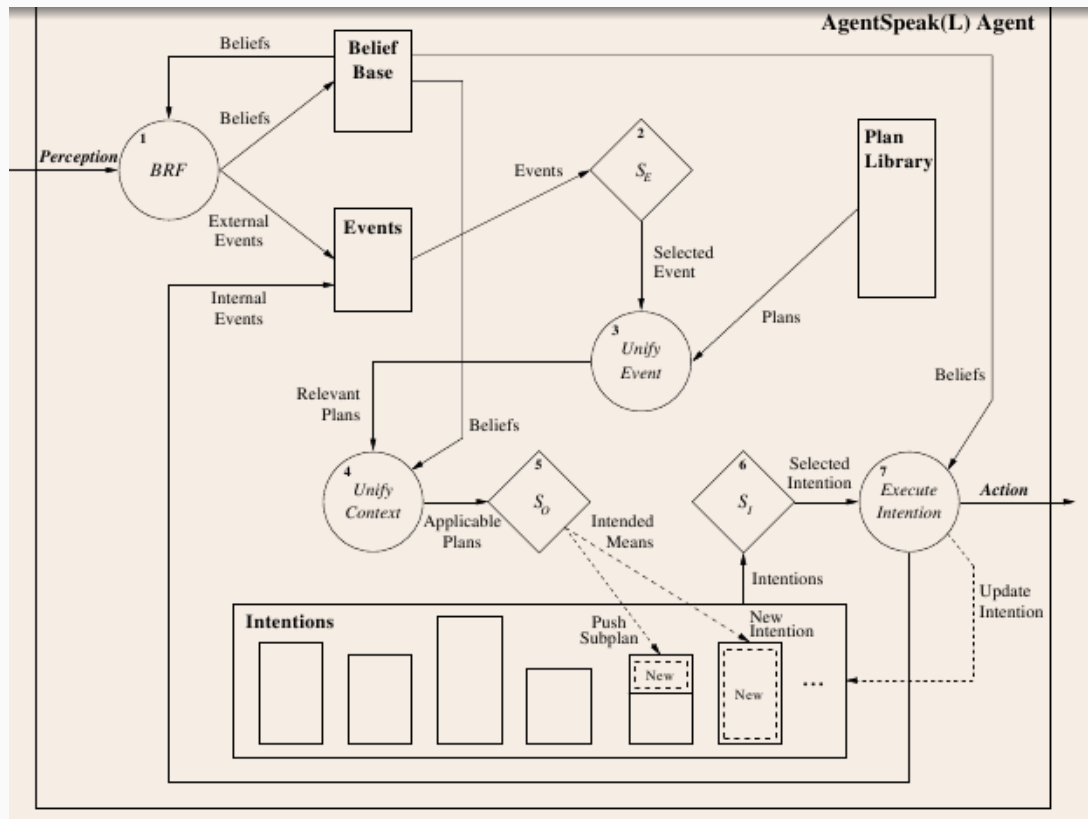
## The BDI Execution Cycle [Rao&Georgeff 92]



### The BDI Execution Cycle: Detailed Version

- 1 **Observe** the environment for new *external* events.
- 2 Pick a pending **event** *e*.
- 3 Select **relevant** plans from library (match event).
- 4 Select **applicable** plans from relevant set (match context).
- 5 If event *e* is **external**, create new intention with selected plan.
- 6 If event *e* is **internal**, update intention with selected plan on top.
- 7 Partially **execute** some intention (may post internal events).
  - ▶ If execution fails, then perform **failure recovery**.
- 8 Repeat cycle.

# Detailed Architecture of AgentSpeak



## The CAN Language [Winikioff et al. 2002]

### CAN: Conceptual Agent Notation

Can be seen as an extension of Rao's AgentSpeak.

A CAN agent is defined as  $Agt = \langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle$ , where:

- ▶  $\mathcal{N}$  is the **agent name**.
- ▶  $\mathcal{B}$  is the **belief base**: current agent's knowledge.
- ▶  $\mathcal{A}$  is the sequence of actions executed so far.
- ▶  $\Pi$  is a **plan library** containing plan rules  $e : \psi \leftarrow P$ :
  - ▶  $e$  is the **triggering event**
  - ▶  $\psi$  is the **context condition**
  - ▶  $P$  is the **plan-body**
- ▶  $\Gamma$  is the **intention base**: partially uninstantiated plan-bodies.

# The CAN Language: Beliefs, Actions & Goals

- ▶ In principle,  $\mathcal{B}$  is any KR formalism that allows queries and updates:
  - ▶  $\mathcal{B} \models \phi$ ,  $\mathcal{B} \cup \{b\}$  and  $\mathcal{B} \setminus \{b\}$ .
- ▶ In practice: a database of facts.
- ▶ If  $b$  is a predicate symbol, and  $t_1, \dots, t_n$  are (first-order) terms,  $b(t_1, \dots, t_n)$  is a **belief atom**.
  - ▶ Ground belief atoms are base beliefs
  - ▶ If  $\phi$  is a belief atom,  $\phi$  and  $\neg\phi$  are **belief literals**.
- ▶ If  $a$  is an action symbol and  $t_1, \dots, t_n$  are first-order terms, then  $a(t_1, \dots, t_n)$  is an **action**.
- ▶ If  $g$  is a predicate symbol, and  $t_1, \dots, t_n$  are terms,  $!g(t_1, \dots, t_n)$  and  $?g(t_1, \dots, t_n)$  are goals
  - ▶ '!' denotes **posting of achievement goals**.
  - ▶ '?' denotes **test goals**.

# The CAN Language: Plans & Intentions

- ▶  $\Pi$  is a **plan library** containing plan rules  $e : \psi \leftarrow P$ :
 

$act$	<i>primitive action</i>
$+b$	<i>belief addition</i>
$-b$	<i>belief deletion</i>
$? \phi$	<i>tests goal</i>
$!e$	<i>achievement event goal</i>
$P_1; P_2$	<i>sequence</i>
$P_1    P_2$	<i>interleaved concurrency</i>

Plus, the following system-auxiliary constructs:

$nil(\theta)$	<i>empty program with bindings</i>
$P_1 \triangleright P_2$	<i>try <math>P_1</math>; else <math>P_2</math></i>
$(\psi_1 : P_1, \dots, \psi_n : P_n)$	<i>guarded plans</i>

- ▶  $\Gamma$  is the **intention base**: set of partially uninstantiated plan-bodies.
  - ▶ E.g.:  $(?phone(john, N); call(N); !talk) || !cook\_dinner$

# Semantics of CAN

CAN has an single-step **operational semantics** (Plotkin 1981):

- ▶ Give **meaning** to computer programs in a mathematically rigorous way.
- ▶ System is interpreted as **sequences of computational steps**. These sequences then are the meaning of the program.
- ▶ Set of rules defining the **transitions** between system **configurations**.
- ▶ Contrast with denotational semantics & axiomatic semantics.

We will use two types of configurations:

- 1 Agent configuration:  $\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle$ .
  - ▶  $\Gamma$  is a set of partially-instantiated plan bodies – the **intention base**.
- 2 Intention configuration:  $\langle \Pi, \mathcal{B}, \mathcal{A}, P \rangle$ .
  - ▶  $P$  is just one partially-instantiated plan body – the **selected intention**.

What we need are rules to state how configurations may evolve (one step).

# Semantics of CAN (cont.)

The semantics of CAN is modularly defined in two levels:

- 1 Agent-level semantics:  $\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', \Gamma' \rangle$ .
  - ▶ State that agent configuration  $\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle$  **may legally evolve** to configuration  $\langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', \Gamma' \rangle$ .
- 2 Intention-level semantics:  $\langle \Pi, \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \langle \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$ .
  - ▶ State that intention configuration  $\langle \Pi, \mathcal{B}, \mathcal{A}, P \rangle$  **may legally evolve** to configuration  $\langle \Pi, \mathcal{B}', \mathcal{A}', P' \rangle$ .

Legal transitions are characterized by a set of rules of the form:

$$\frac{\text{Set of conditions}}{C \longrightarrow C'} \text{ RuleName}$$

## Definition (BDI Agent Execution)

A BDI execution  $E$  of an agent  $C_0 = \langle \mathcal{N}, \Pi, \mathcal{B}_0, \mathcal{A}_0, \Gamma_0 \rangle$  is a, possibly infinite, sequence of agent configurations  $C_0 \cdot C_1 \cdot \dots$  such that  $C_i \Longrightarrow C_{i+1}$ , for every  $i \geq 0$ . A terminating execution is a finite execution  $C_0 \cdot \dots \cdot C_n$  with  $\Gamma_n = \{\}$ .

# Agent-Level Semantics

$$\frac{P \in \Gamma \quad \langle \mathcal{B}, \mathcal{A}, P \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}', \mathcal{A}', (\Gamma \setminus \{P\}) \cup \{P'\} \rangle} \text{Agt}_{step}$$

$$\frac{e \text{ is a new external event}}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \cup \{!e\} \rangle} \text{Agt}_{event}$$

$$\frac{P \in \Gamma \quad \langle \mathcal{B}, \mathcal{A}, P \rangle \not\longrightarrow}{\langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \rangle \Longrightarrow \langle \mathcal{N}, \Pi, \mathcal{B}, \mathcal{A}, \Gamma \setminus \{P\} \rangle} \text{Agt}_{clean}$$

**Execute an active intention  $P$ .   Assimilate an external event  $e$ .**  
**Remove an active intention  $P$  that is blocked.**

# Intention-Level Semantics: Basic Programs

$$\frac{\mathcal{B} \models \phi\theta}{\langle \mathcal{B}, \mathcal{A}, ?\phi \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, nil(\theta) \rangle} ?$$

$$\overline{\langle \mathcal{B}, \mathcal{A}, act \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A} \cdot act, nil(\emptyset) \rangle} \text{do}$$

$$\overline{\langle \mathcal{B}, \mathcal{A}, +b \rangle \longrightarrow \langle \mathcal{B} \cup \{b\}, \mathcal{A}, nil(\emptyset) \rangle} +b$$

$$\overline{\langle \mathcal{B}, \mathcal{A}, -b \rangle \longrightarrow \langle \mathcal{B} \setminus \{b\}, \mathcal{A}, nil(\emptyset) \rangle} -b$$

**Goal test condition – propagate corresponding bindings.**  
**Primitive action execution – actions just execute.**  
**Addition & deletion of a belief atom.**

# Intention-Level Semantics: Complex Programs

$$\frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_1 \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1; P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_1; P_2 \rangle} \text{Seq} \quad \frac{}{\langle \mathcal{B}, \mathcal{A}, \text{nil}(\theta) ; P \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, P\theta \rangle} \text{Seq}_t$$

$$\frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \parallel P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \parallel P_2 \rangle} \parallel_1 \quad \frac{}{\langle \mathcal{B}, \mathcal{A}, \text{nil} \parallel P_2 \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, P_2 \rangle} \parallel_{t_1}$$

$$\frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \triangleright P_2 \rangle} \triangleright \quad \frac{}{\langle \mathcal{B}, \mathcal{A}, \text{nil}(\theta) \triangleright P' \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, \text{nil}(\theta) \rangle} \triangleright_t$$

$$\frac{P_1 \neq \text{nil} \quad \langle \mathcal{B}, \mathcal{A}, P_1 \rangle \not\longrightarrow \quad \langle \mathcal{B}, \mathcal{A}, P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle} \triangleright_f$$

Sequence of programs – propagate corresponding bindings.

Interleaved concurrency – two analogous rules for  $P_2$ .

Try execution – jump to  $P_2$  if  $P_1$  is not working.

# Intention-Level Semantics: Selection & Failure

When an unresolved event is addressed:

- 1 inspect plan-library for potential *relevant* plans;
- 2 select one of those plans that is *applicable*;
- 3 start executing.

When the plan  $P$  being pursued for an event  $e$  has problems executing:

- 1 is there an *alternative applicable* strategy  $P'$  we can follow?
  - maybe same plan-strategy but with different variable bindings.
- 2 if so, switch to it!
- 3 otherwise, go up in the hierarchy of goals and do the same reasoning.

All this is achieved by means of the following two special constructs:

$P_1 \triangleright P_2$   
 $(\psi_1 : P_1, \dots, \psi_n : P_n)$

*try  $P_1$  if possible; else  $P_2$*   
*guarded (relevant) plans*

# Intention-Level Semantics: Selection & Failure (cont.)

$$!e \longrightarrow (\psi_1 : P_1, \dots, \psi_n : P_n) \longrightarrow$$

$$P_i \theta_i \triangleright (\psi_1 : P_1, \dots, \underline{\psi_i \wedge \vec{x} \neq \theta_i} : P_i, \dots, \psi_n : P_n) \xrightarrow{*}$$

$$nil \triangleright (\psi_1 : P_1, \dots, \psi_i \wedge \vec{x} \neq \theta_i : P_i, \dots, \psi_n : P_n) \longrightarrow nil$$

$$P'_i \theta_i \triangleright (\psi_1 : P_1, \dots, \psi_i \wedge \vec{x} \neq \theta_i : P_i, \dots, \psi_n : P_n) \longrightarrow$$


$$P_j \theta_j \triangleright (\psi_1 : P_1, \dots, \psi_i \wedge \vec{x} \neq \theta_i : P_i, \dots, \underline{\psi_j \wedge \vec{x} \neq \theta_j} : P_j, \dots, \psi_n : P_n)$$

$$\frac{\Delta = \{\psi_i \theta : P_i \theta \mid e' : \psi_i \leftarrow P_i \in \Pi \wedge \theta = \text{mgu}(e, e')\}}{\langle \mathcal{B}, \mathcal{A}, !e \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, \langle \Delta \rangle \rangle} \text{Event}$$

$$\frac{\psi_i(\vec{x}) : P_i \in \Delta \quad \mathcal{B} \models \psi_i(\vec{x}) \theta}{\langle \mathcal{B}, \mathcal{A}, \langle \Delta \rangle \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, P_i \theta \triangleright ((\Delta \setminus \{\psi_i(\vec{x}) : P_i\}) \cup \{\psi_i(\vec{x}) \wedge \vec{x} \neq \theta : P_i\}) \rangle} \text{Sel}$$

$$\frac{\langle \mathcal{B}, \mathcal{A}, P_1 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P' \triangleright P_2 \rangle} \triangleright_t \frac{\langle \mathcal{B}, \mathcal{A}, (nil \triangleright P_2) \rangle \longrightarrow \langle \mathcal{B}, \mathcal{A}, nil \rangle}{\langle \mathcal{B}, \mathcal{A}, P_1 \triangleright P_2 \rangle \longrightarrow \langle \mathcal{B}', \mathcal{A}', P'_2 \rangle} \triangleright_f$$

# Failure at Work within the Goal Hierarchy





# Some Other Formal BDI Languages

## AgentSpeak

[LNCS 1996]

- ▶ no concurrency in plans; no built-in failure handling.
- ▶ special belief updates events  $! + b$  and  $! - b$ .

## 3APL

[Autonomous Agents and Multi-Agent Systems 1999]

- ▶ no events, but uses a goal base  $\mathcal{G}$ ;
- ▶ different type of (practical) rules in priority: failure rules, reactive rules, plan rules, optimization rules.

$$\pi_h \leftarrow \varphi \mid \pi_b$$

## GOAL

[Journal of Applied Logic 2007]

- ▶ no events, but uses goal base  $\Pi$ .
- ▶ two special goal adoption actions:  $\text{adopt}(\phi)$  and  $\text{drop}(\phi)$ .
- ▶ plan rules are called *conditional actions*:  $\varphi \triangleright \text{do}(a)$ .

$$\mathbf{B}(\text{in\_cart}(\text{book}) \wedge \mathbf{G}(\text{bought}(\text{book}))) \triangleright \text{do}(\text{pay\_cart})$$

# Review

In this lecture we have seen:

## 1 Basic concepts of BDI programming:

- ▶ programming using mentalistic concepts such as beliefs, desires, capabilities, etc.
- ▶ goal-oriented programming – via **events**;
- ▶ implicit programming – via **plan library & context conditions**;
- ▶ rational execution cycle: on-the-fly recombination of plans.

## 2 CAN formal BDI programming language:

- ▶ captures the basic notions of BDI programming: rational executor;
- ▶ formal operational semantics;
- ▶ includes built-in failure handling.

## Next Lecture

# Declarative Goals in CAN & Hierarchical planning in CAN

## Next Lecture

In the next lecture we will:

- 1 Show how to accommodate hierarchical HTN-style planning into CAN.
  - ▶ to perform some “offline” look-ahead reasoning within the whole online “reactive” execution scheme.
- 2 Review the Java-based JACK agent programming language:
  - ▶ go over a basic gold-mining agent team implementation.

## BDI Formal Languages



Anand S. Rao and Michael P. Georgeff.

An abstract architecture for rational agents.

In *Proceedings of KR-92*, pages 438–449, 1992.



Koen V. Hindriks, Frank S. de Boer, Wiebe van der Hoek, and John-Jules Ch. Meyer.

Agent programming in 3APL.

*Autonomous Agents and Multi-Agent Systems*, 2:357–401, 1999.



Frank S. de Boer, Koen V. Hindriks, Wiebe van der Hoek, and John-Jules Ch. Meyer.

A Verification Framework for Agent Programming with Dec. Goals.

*Journal of Applied Logic*, 5(2):277–302, 2007.



Michael Winikoff, Lin Padgham, James Harland, and John Thangarajah.

Declarative & procedural goals in intelligent agent systems.

In *Proceedings of the KR-02*, 2002.

## BDI Systems I



Georgeff and Lansky.

Reactive reasoning and planning.

In *Proceedings of AAAI-87*, pages 677–682, 1987.



Georgeff, M. P. and Ingrand, F. F.

Decision-making in an embedded reasoning system.

In *Proceedings of IJCAI-89*, 972–978, 1989.



F. Ingrand, M. Georgeff, and A Rao.

An architecture for real-time reasoning and system control.

*IEEE Expert*, 7(6), 1992.



Mark d’Inverno, Michael Luck, Michael Georgeff, , David Kinny, and Michael Wooldridge.

The dMARS architechure: A specification of the distributed multi-agent reasoning system.

*Autonomous Agents and Multi-Agent Systems*, 9(1–2):5–53, 2004.

# BDI Systems II



Huber, Marcus J.

JAM: A BDI-theoretic mobile agent architecture.

In *Proceedings of AGENTS-99*, pages 236–243, 1999.



P. Busetta, Ralph Rönquist, A. Hodgson, and A. Lucas.

JACK Intelligent Agents: Components for intelligent agents in Java.

AgentLink News Letter, Jan 1999.



R. H. Bordini and J. F. Hubner.

BDI agent programming in AgentSpeak using Jason.

In *Proceedings of the International Workshop on Computational Logic in Multi-Agent Systems (CLIMA)*, pages 143–164, 2006.



Alexander Pokahr, Lars Braubach, Winfried Lamersdorf

Jadex: Implementing a BDI-Infrastructure for JADE Agents.

In *Search of Innovation (Special Issue on JADE)*, 3(5):76–85, Telecom Italia Lab, Turin, Italy, September 2003, 76-85, 2003.